

The background image shows a vast mountain range under a sky transitioning from deep blue to vibrant orange and pink. In the foreground, a bright orange tent is set up on a rocky, grassy slope, its light reflecting softly. A small body of water is visible in the middle ground, reflecting the colors of the sky.

DEELEARNING

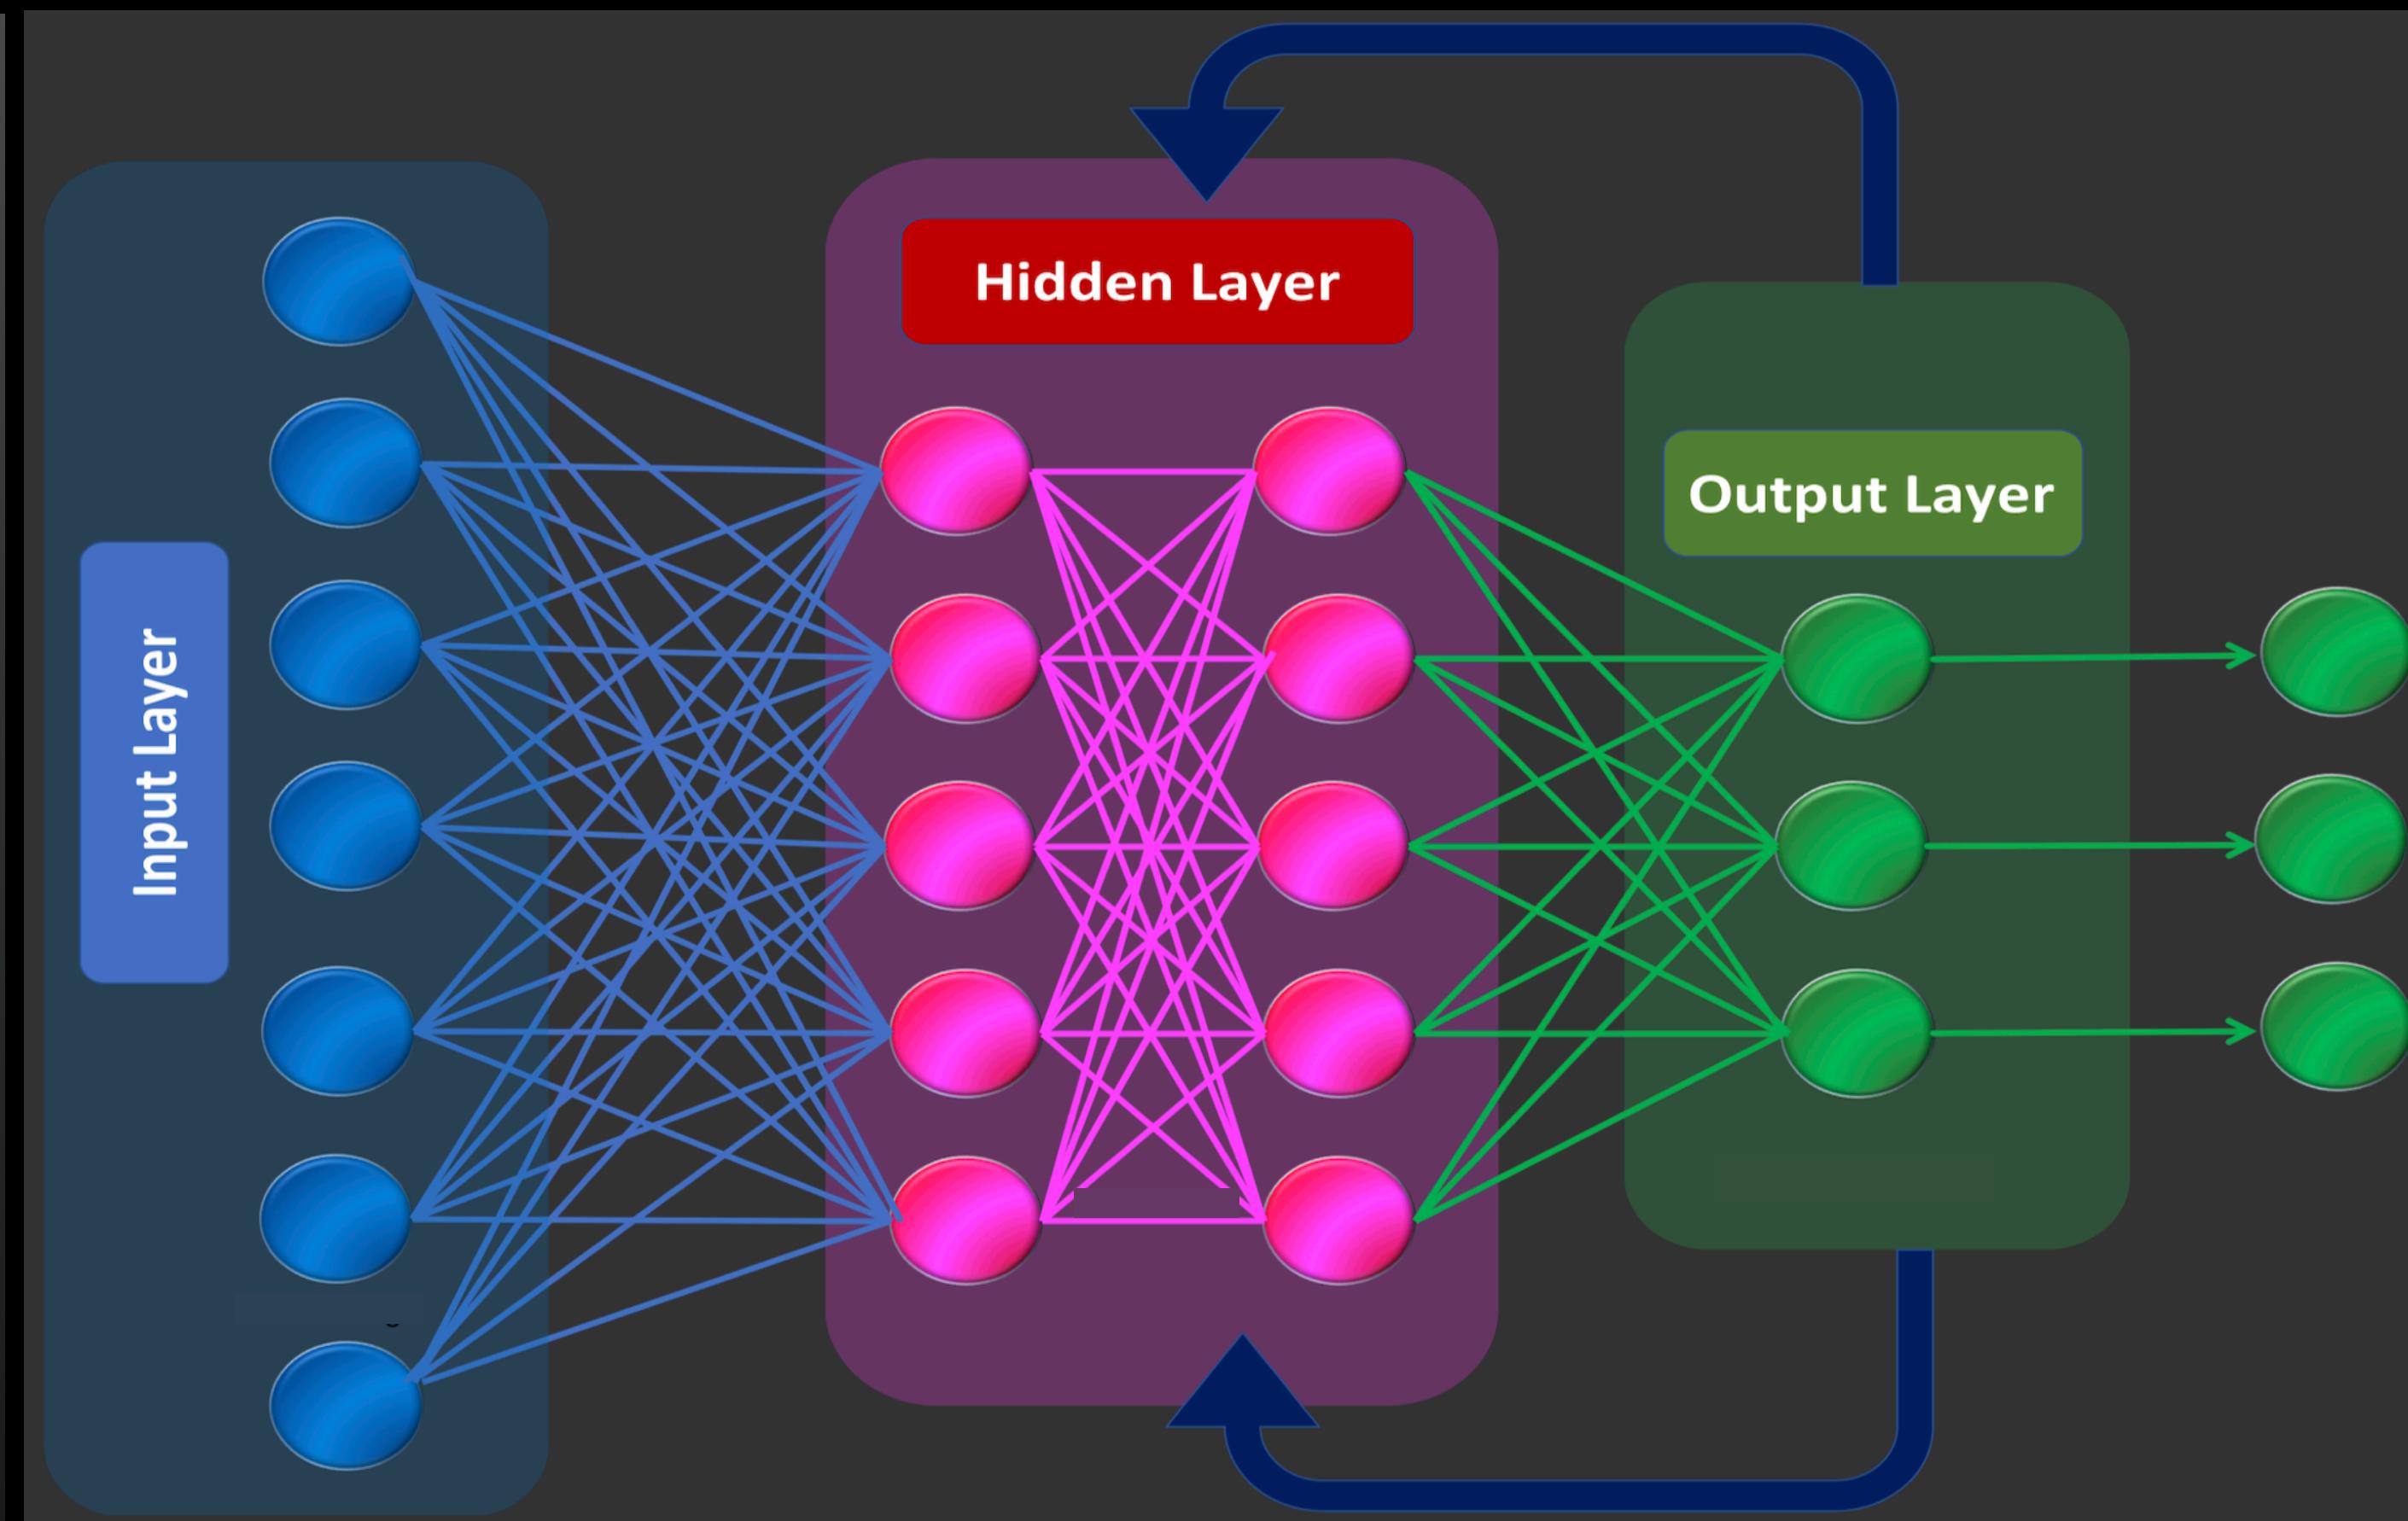
TENSORFLOW

<https://www.tensorflow.org>

Rafael Vieira Coelho

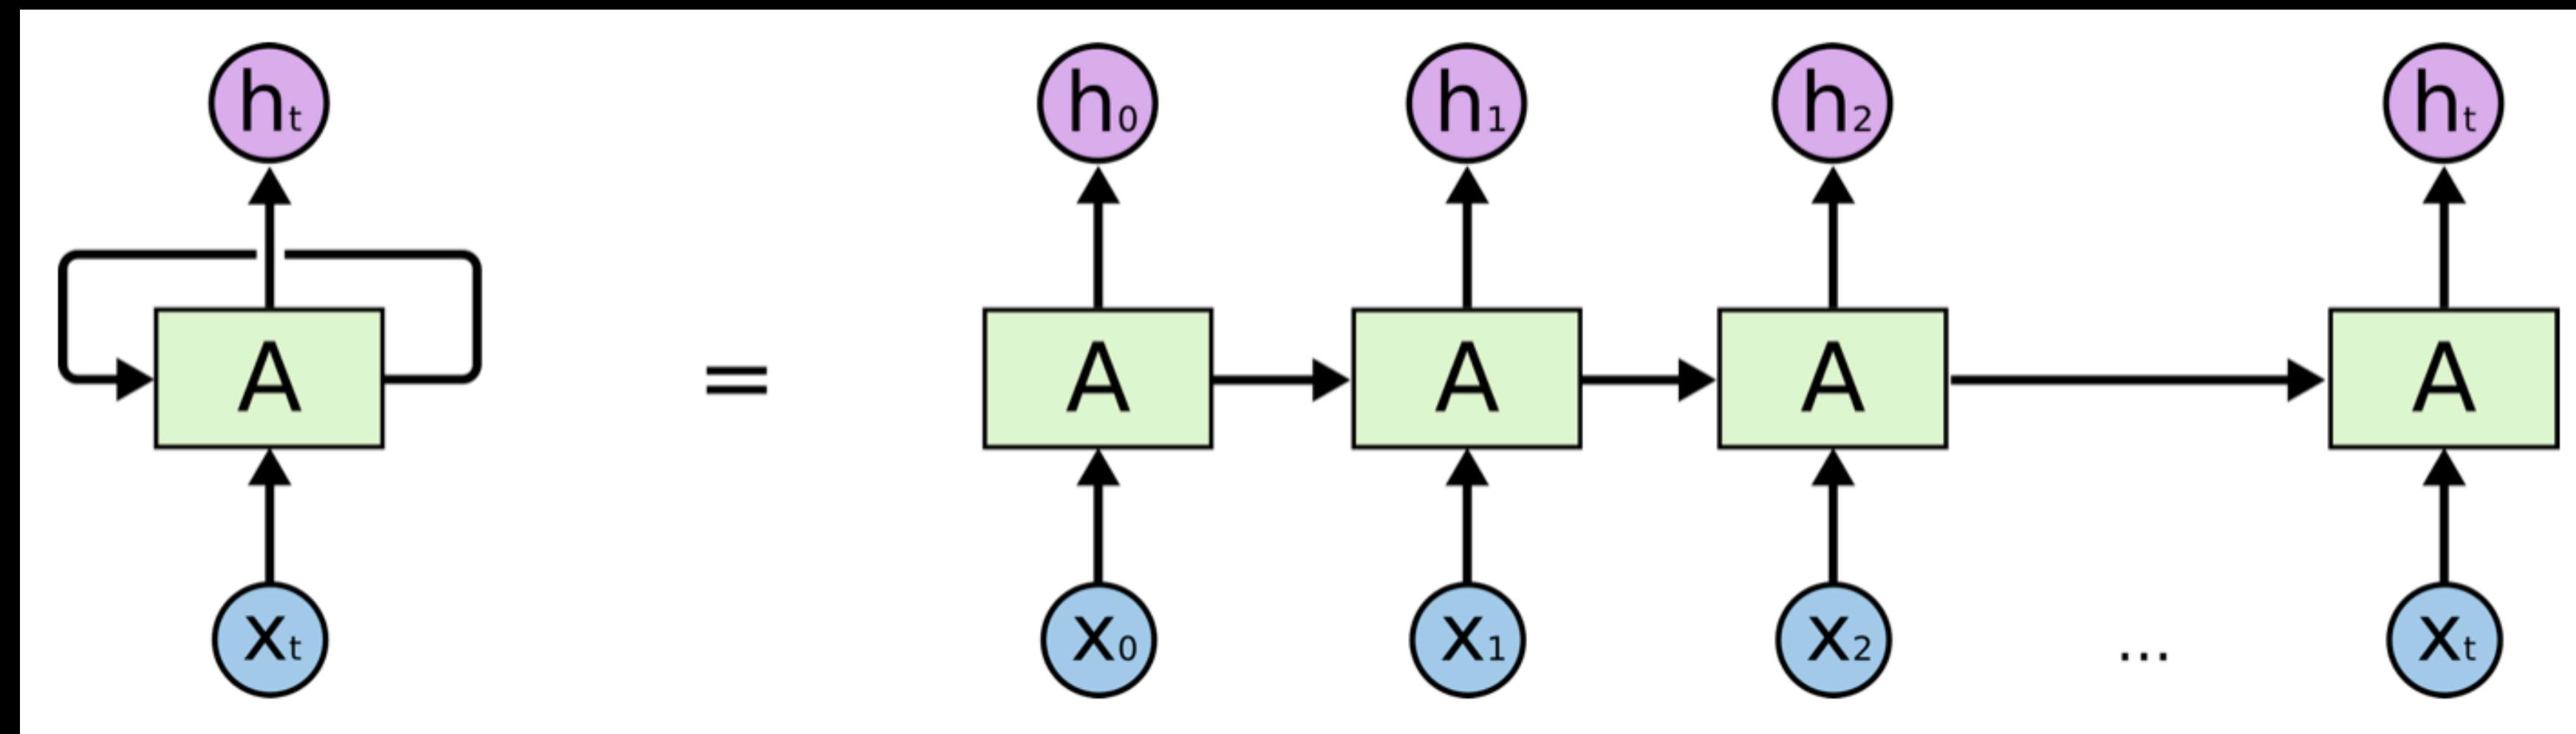
TENSORFLOW

REDES RECORRENTES



A) ARQUITETURA DE REDES NEURAIS LONG SHORT TERM MEMORY (LSTM)

É um tipo de rede neural recorrente, que é usada em diversos cenários de Processamento de Linguagem Natural. Ela pode ser imaginada como múltiplas cópias da mesma rede, cada uma passando uma mensagem a um sucessor.



A LSTM é bem adequada para classificar, processar e prever séries temporais com intervalos de tempo de duração desconhecida.

O QUE É UMA REDE NEURAL LSTM?

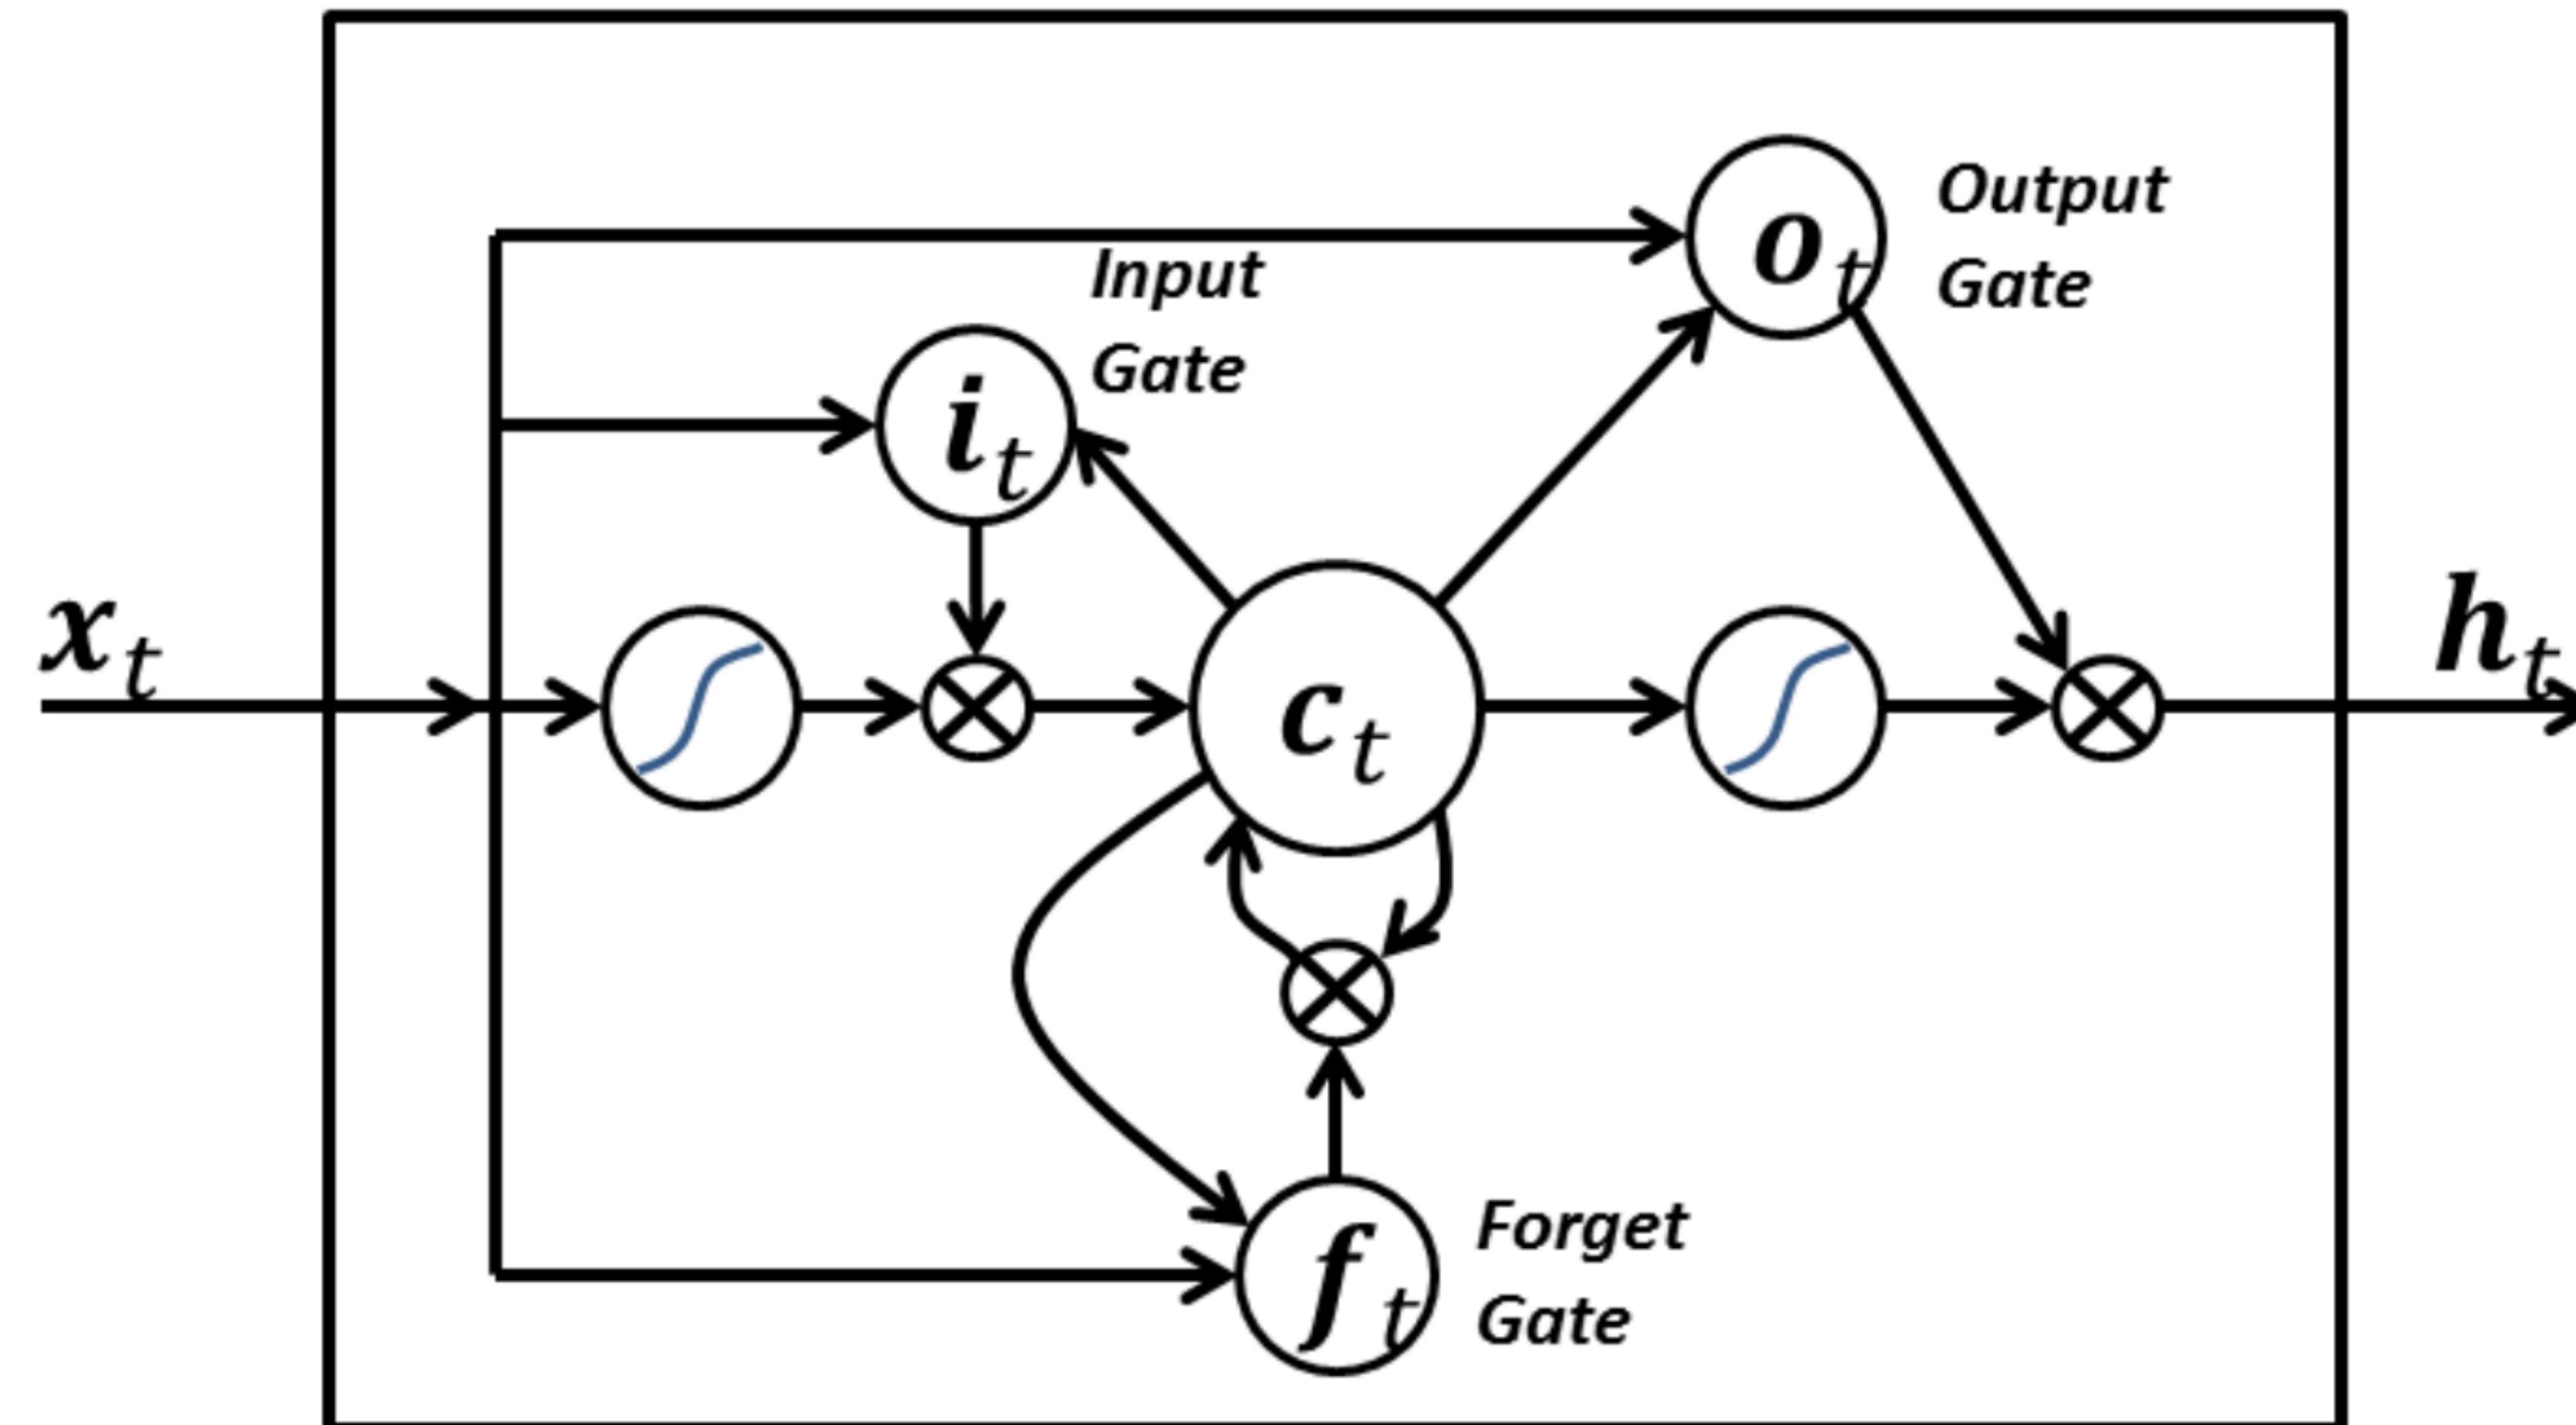
- Os LSTMs contêm informações fora do fluxo normal da rede recorrente em uma célula fechada. As informações podem ser armazenadas, escritas ou lidas a partir de uma célula, como dados na memória de um computador.
- A célula toma decisões sobre o que armazenar, e quando permitir leituras, gravações e exclusões, através de portões abertos e fechados.
- Ao contrário do armazenamento digital em computadores, no entanto, esses portões são analógicos, implementados com a multiplicação de elementos por sigmóides, que estão todos na faixa de 0-1.

COMO FUNCIONAM OS PORTÕES (GATES) EM UMA LSTM?

- Esses portões atuam sobre os sinais que recebem e, de forma semelhante aos nós da rede neural, eles bloqueiam ou transmitem informações com base em sua força e importância, que eles filtram com seus próprios conjuntos de pesos.
- Esses pesos, como os pesos que modulam a entrada e estados ocultos, são ajustados através do processo de aprendizagem das redes recorrentes.
- Ou seja, as células aprendem quando permitir que os dados entrem, saiam ou sejam excluídos através do processo iterativo de fazer suposições, calculando o erro durante o backpropagation e ajustando pesos através da descida do gradiente.

COMO FUNCIONAM OS PORTÕES (GATES) EM UMA LSTM?

O diagrama abaixo ilustra como os dados fluem através de uma célula de memória e são controlados por seus portões.



LSTM: CAMADA ÚNICA

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(tokenizer.vocab_size, 64),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64)),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, None, 64)	523840
bidirectional_1 (Bidirection)	(None, 128)	66048
dense_4 (Dense)	(None, 64)	8256
dense_5 (Dense)	(None, 1)	65
Total params:	598,209	
Trainable params:	598,209	
Non-trainable params:	0	

Long Short Term Memory
(LSTM)

LSTM: DUAS CAMADAS

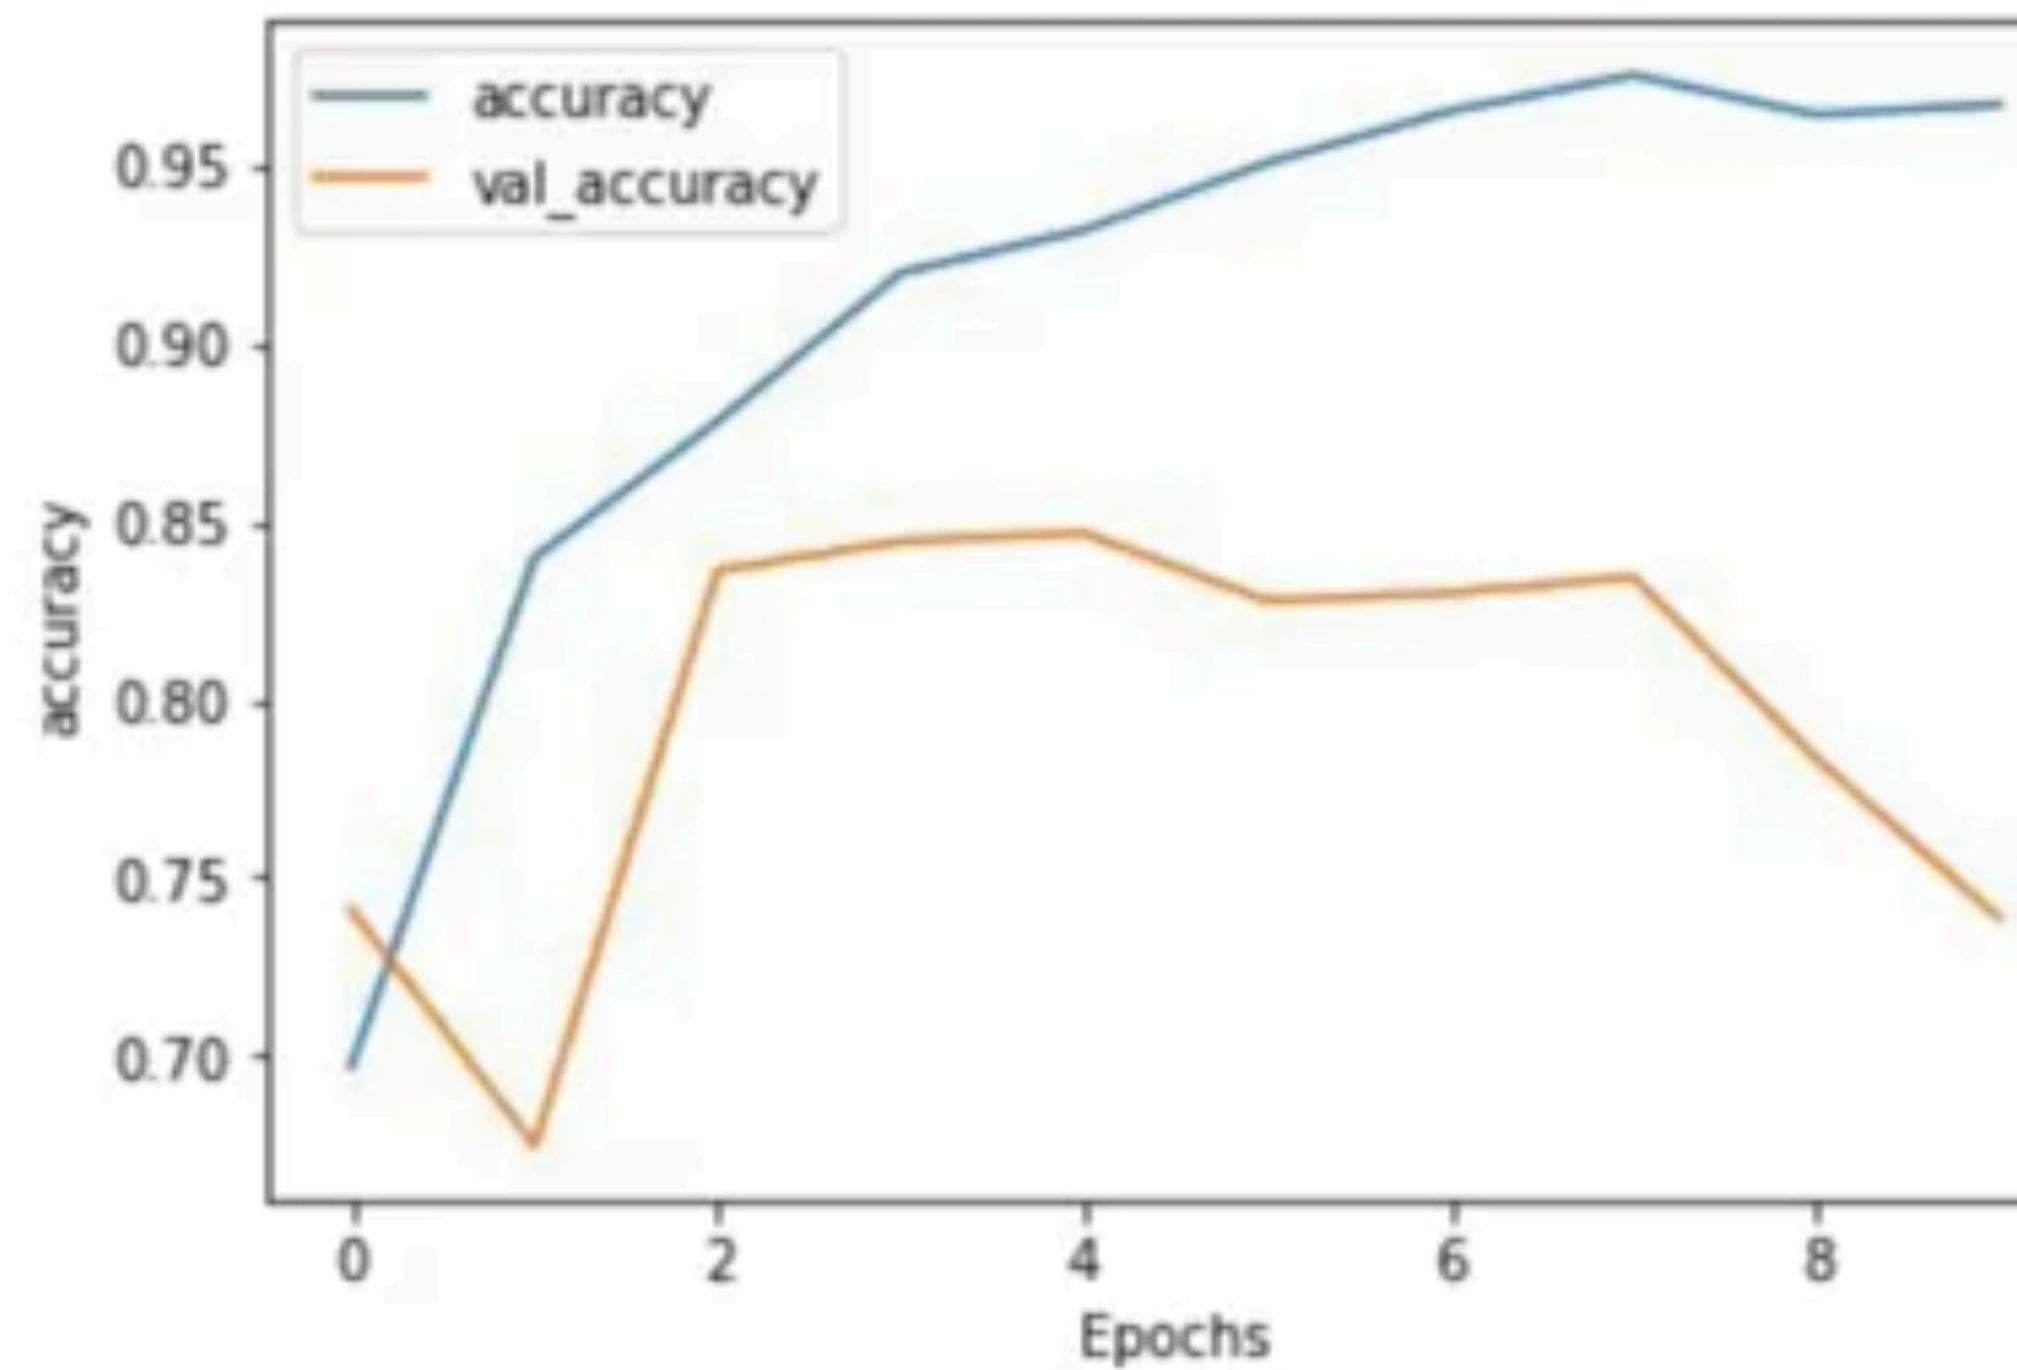
```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(tokenizer.vocab_size, 64),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64, return_sequences=True)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

É necessário estar com True a propriedade `return_sequences` para passar para a próxima camada bidirecional.

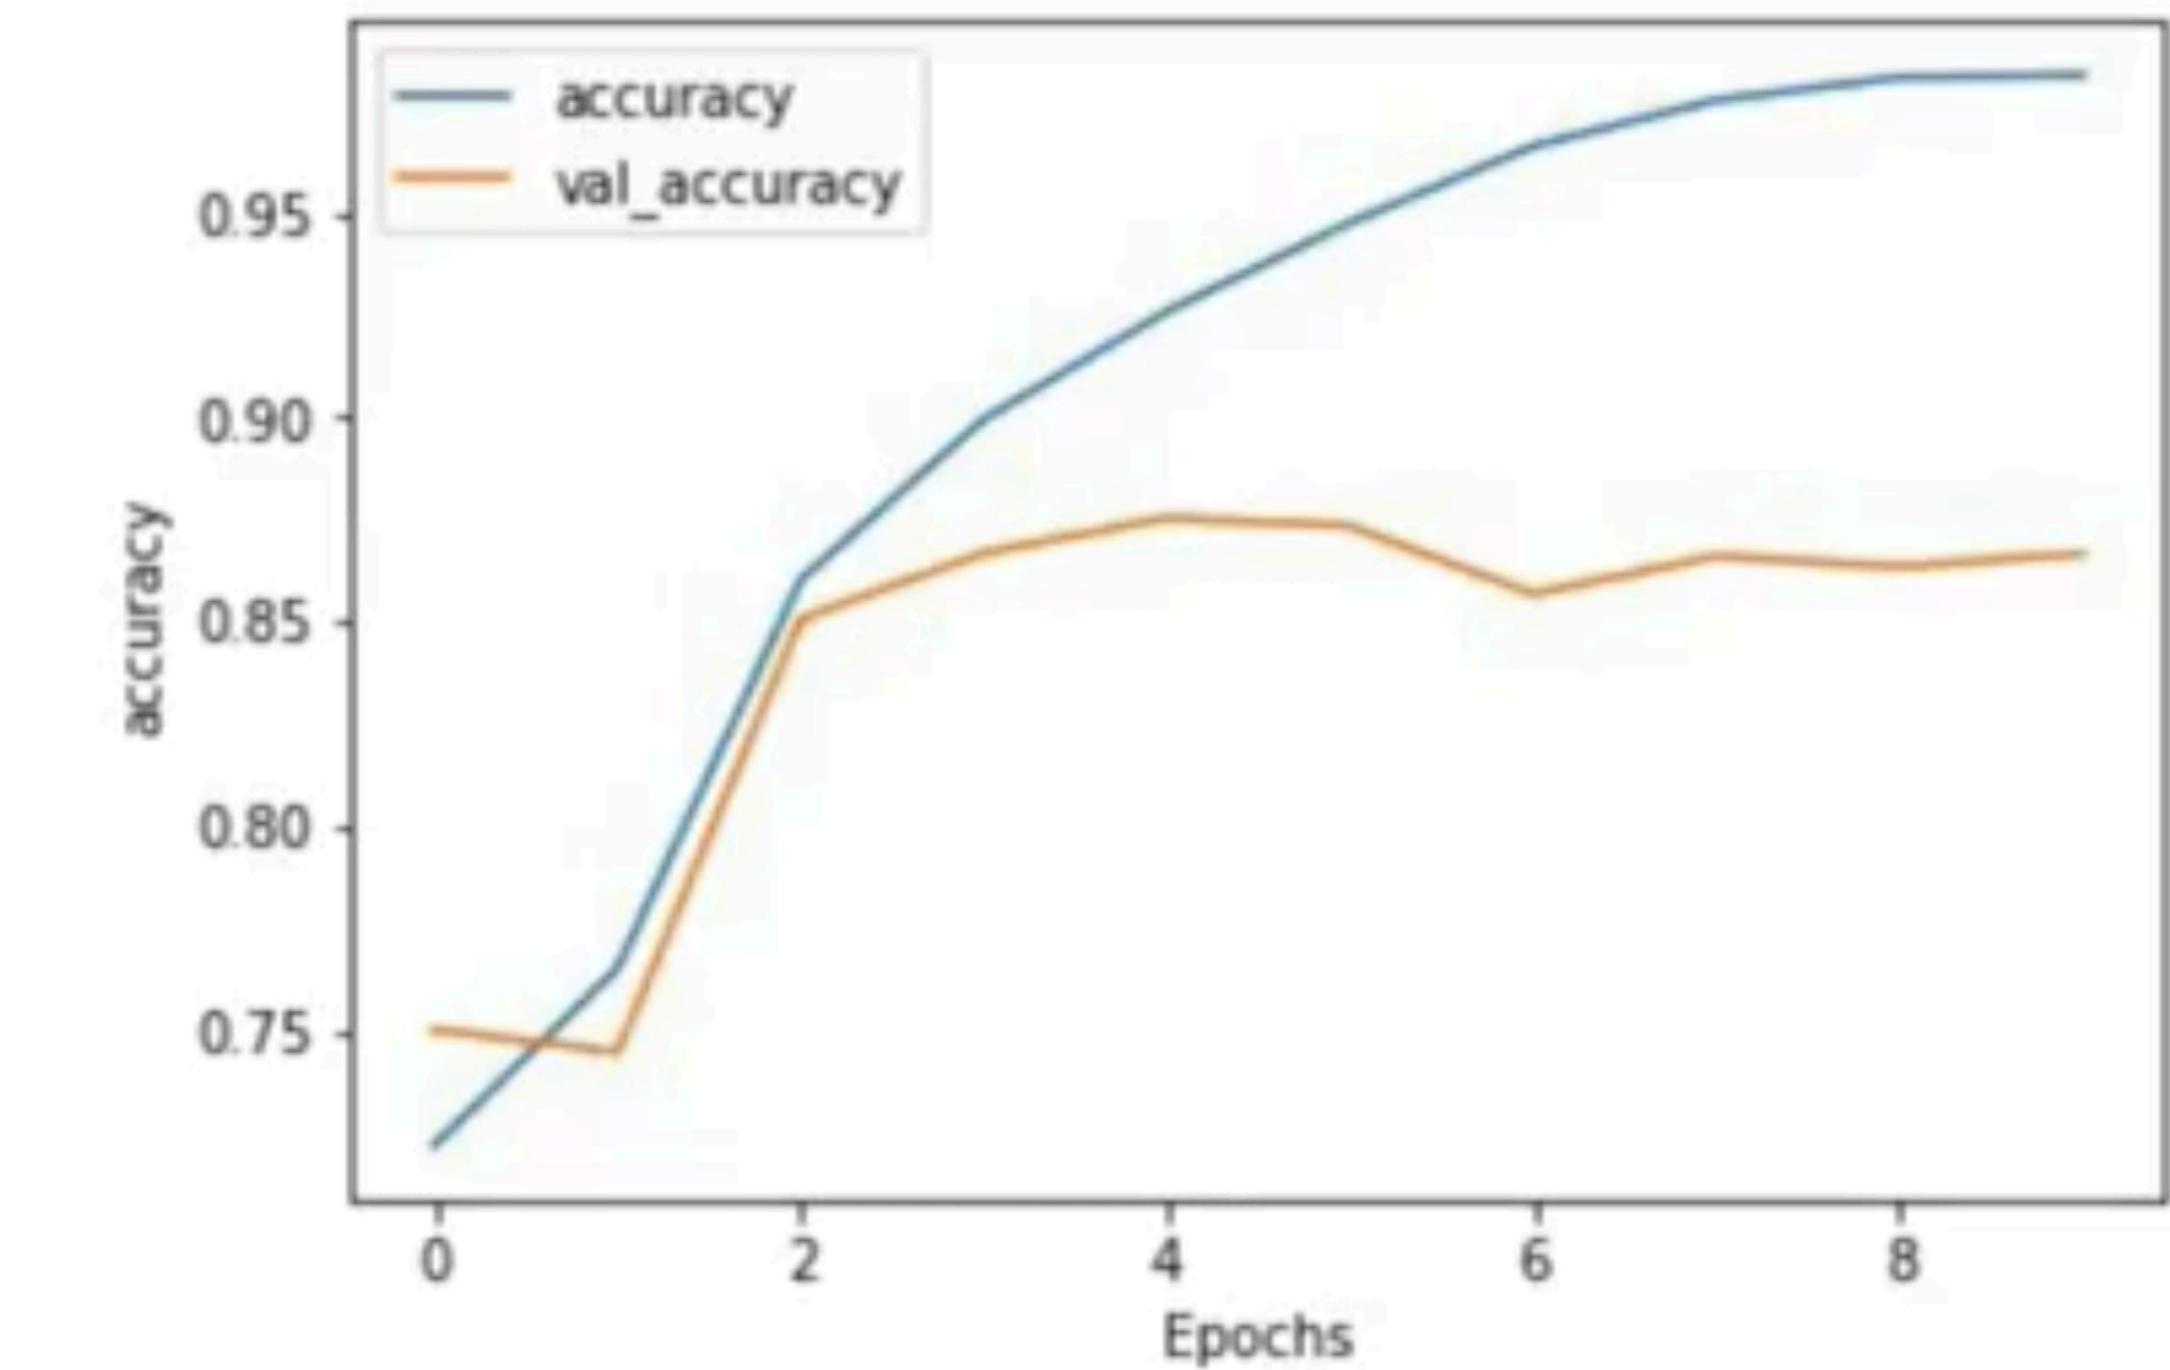
Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, None, 64)	523840
bidirectional_2 (Bidirection)	(None, None, 128)	66048
bidirectional_3 (Bidirection)	(None, 64)	41216
dense_6 (Dense)	(None, 64)	4160
dense_7 (Dense)	(None, 1)	65
=====		
Total params: 635,329		
Trainable params: 635,329		
Non-trainable params: 0		

LSTM: CAMADA ÚNICA X DUAS CAMADAS

10 Epochs : Accuracy Measurement



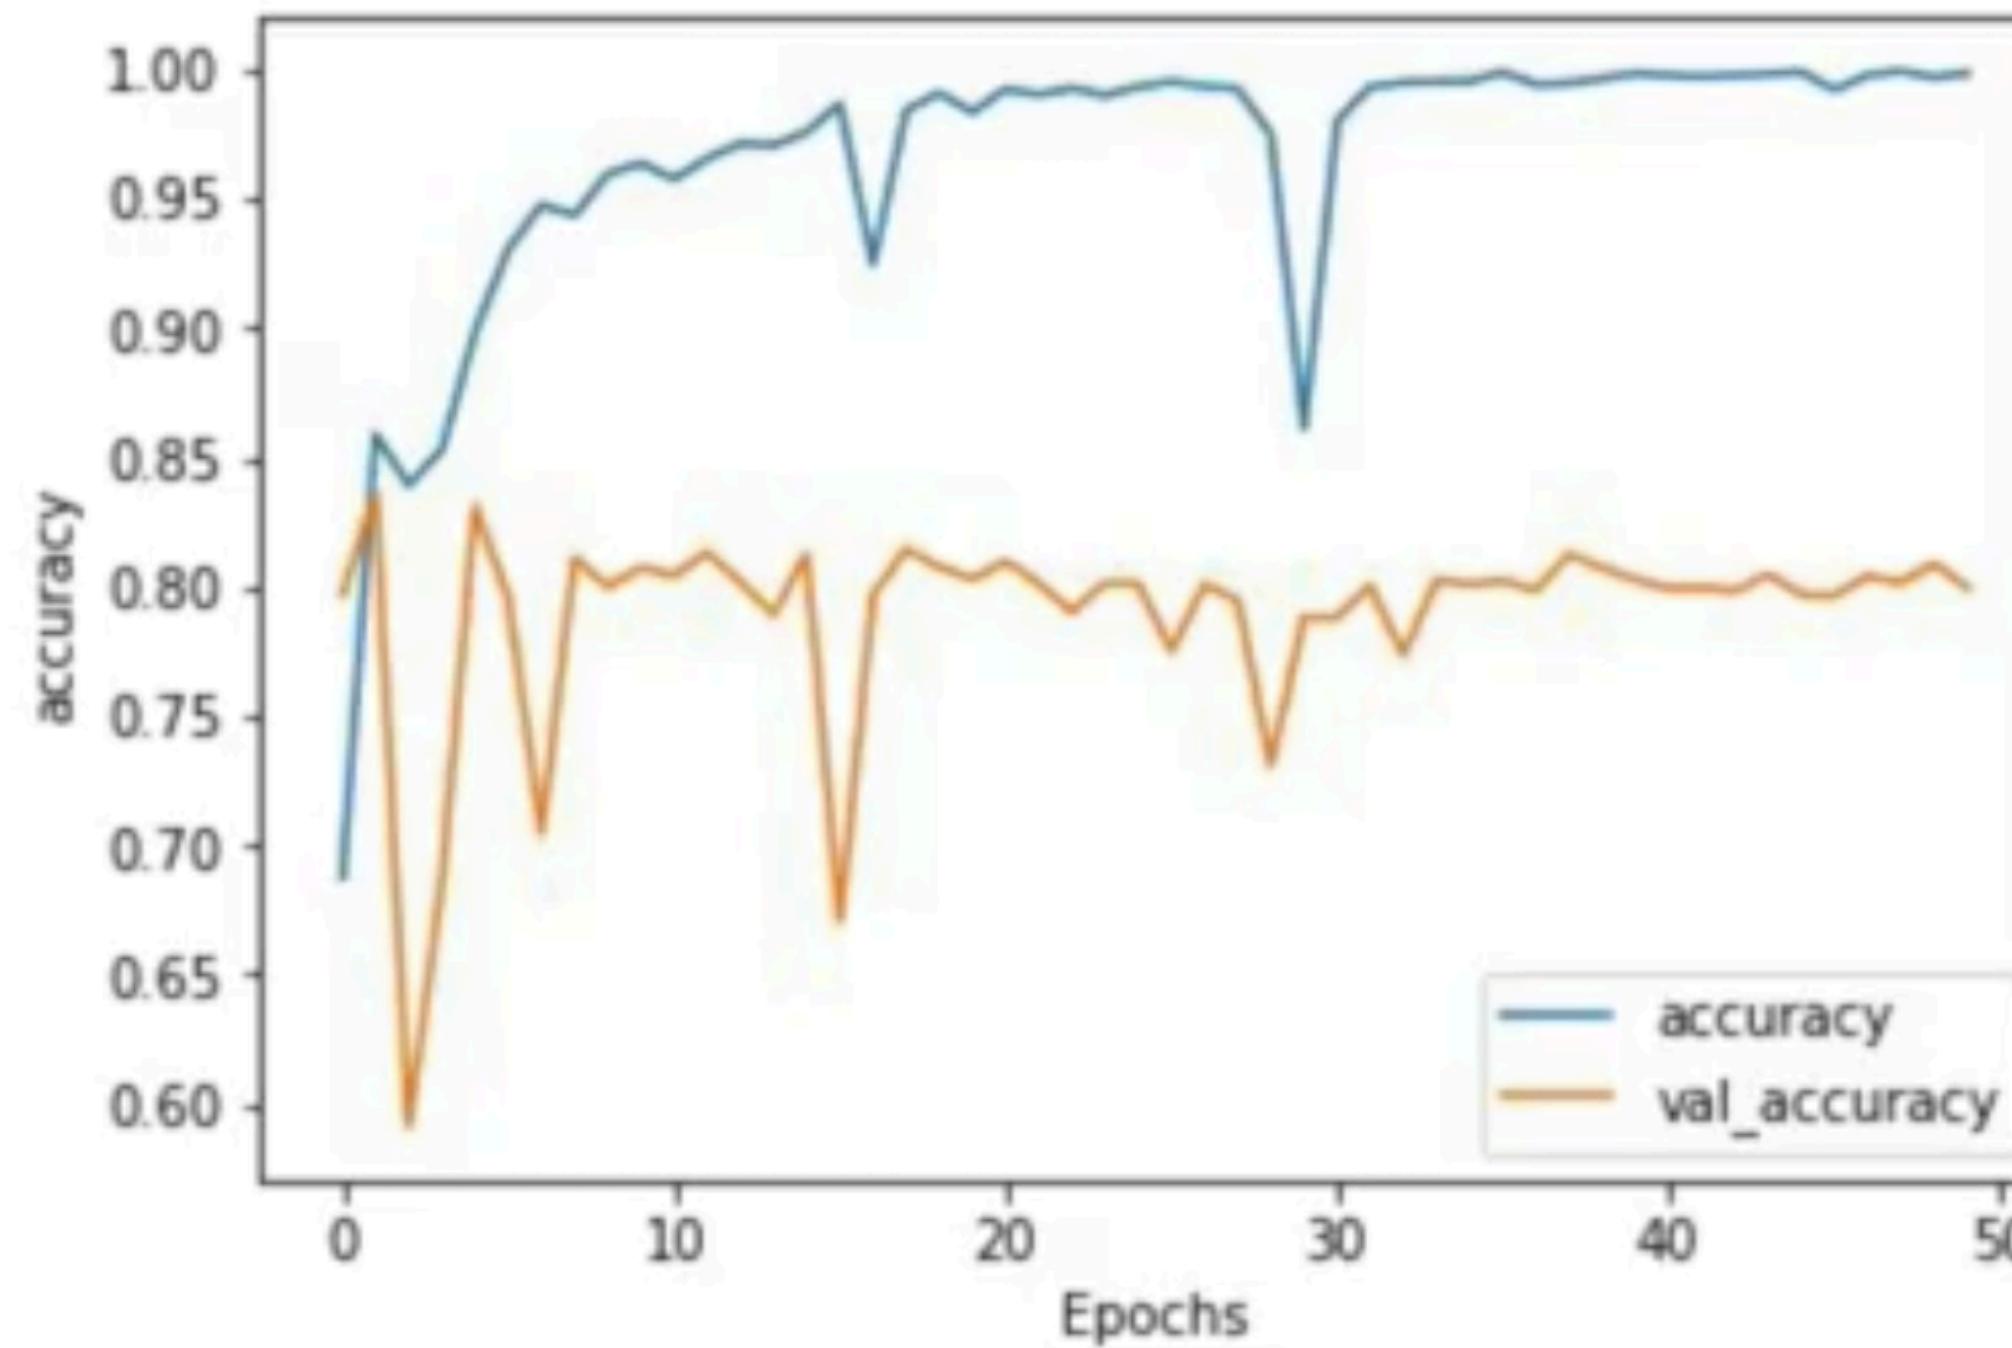
1 Layer LSTM



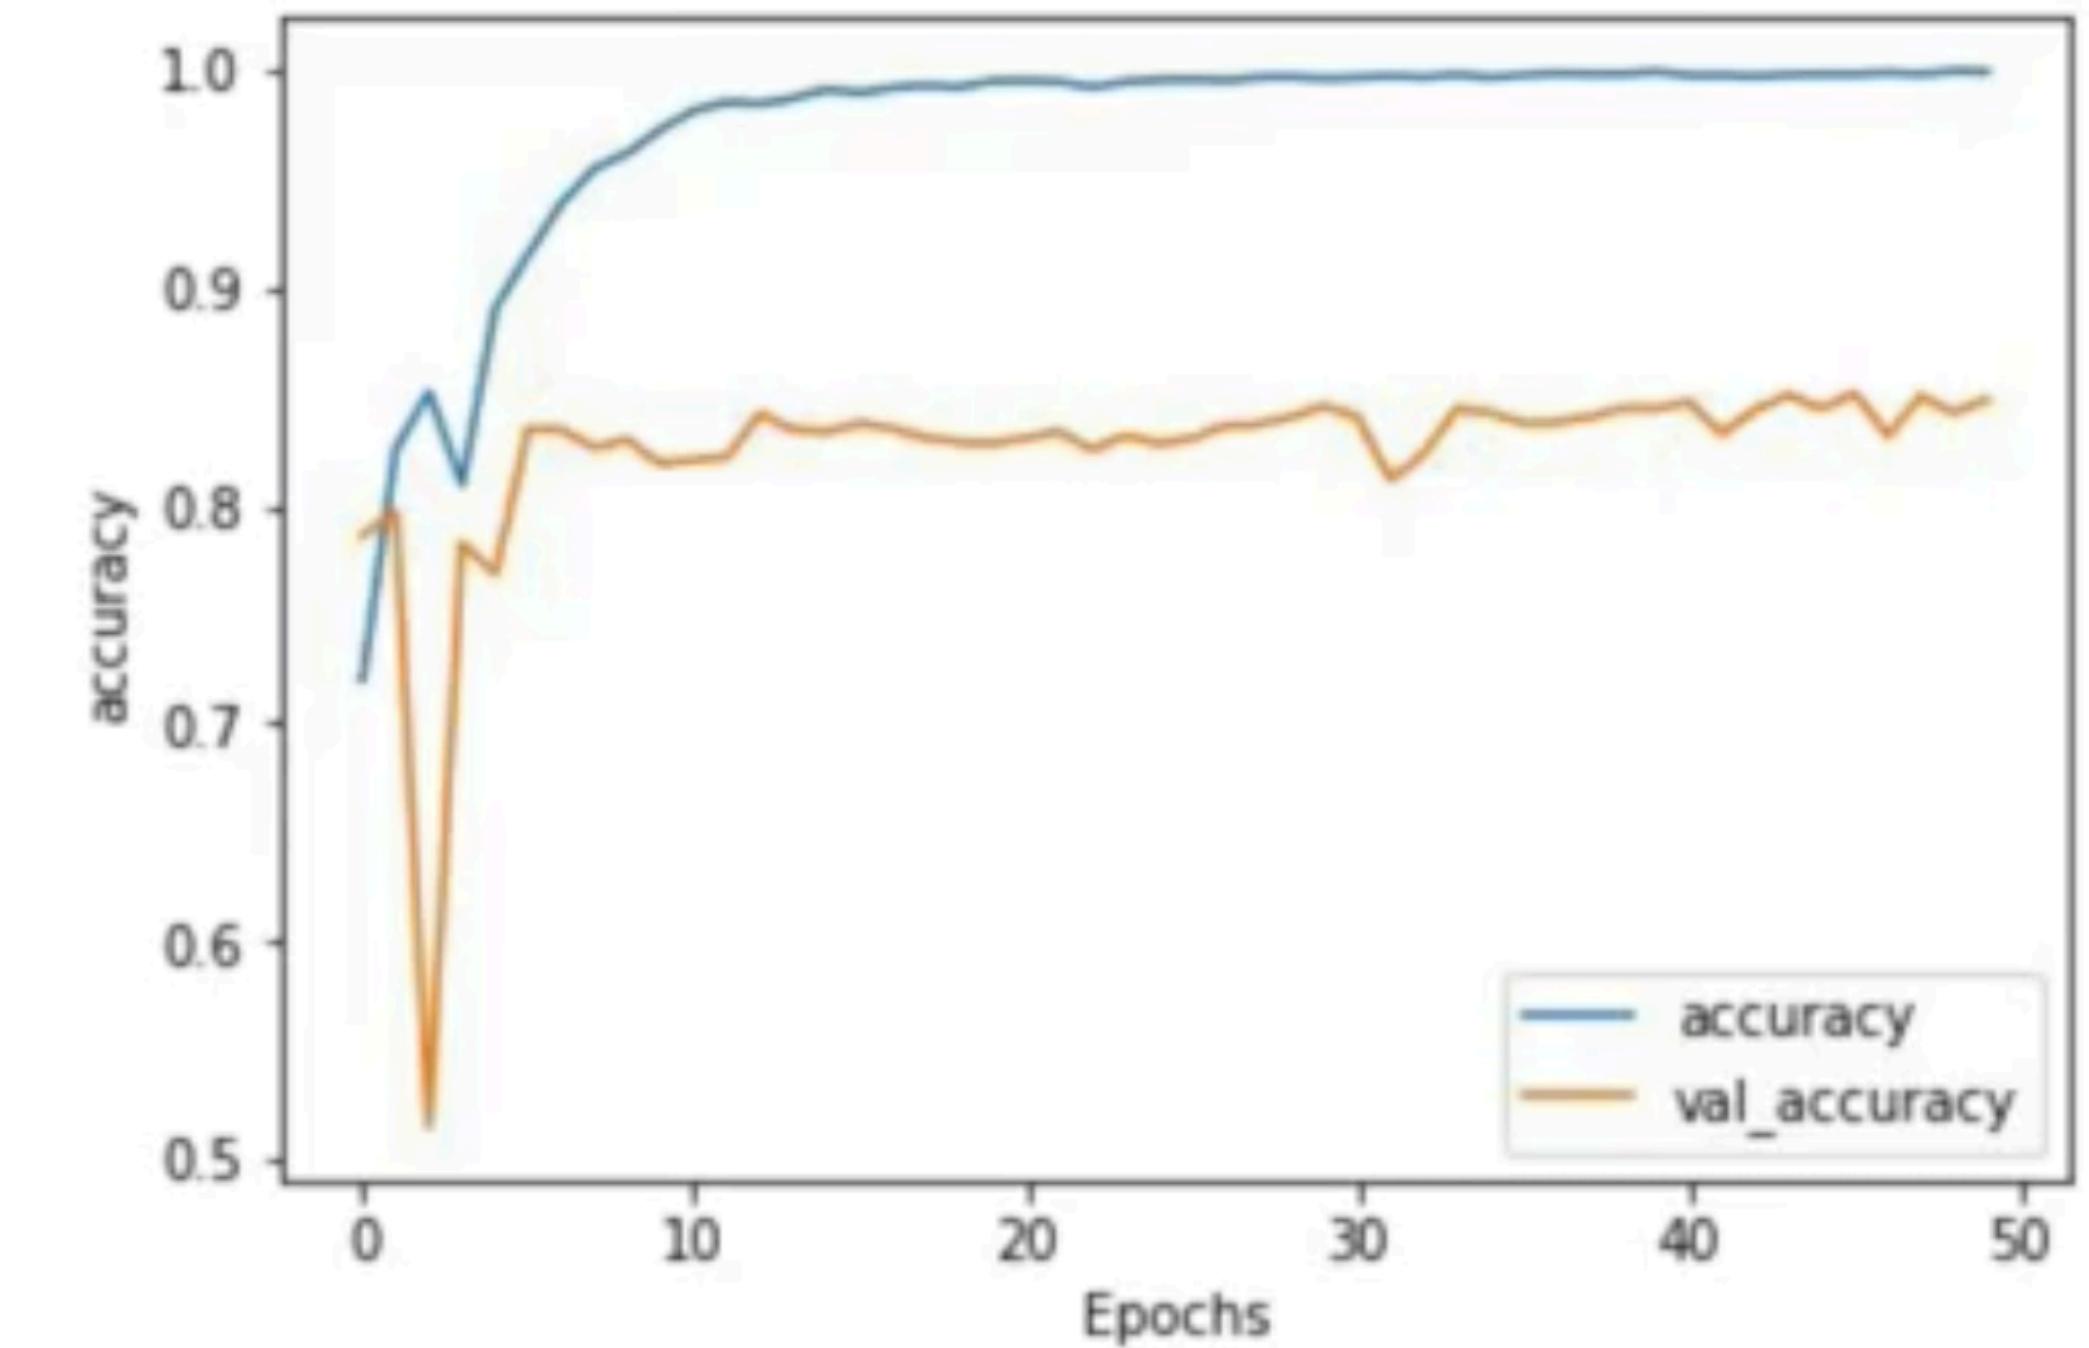
2 Layer LSTM

LSTM: CAMADA ÚNICA X DUAS CAMADAS

50 Epochs : Accuracy Measurement



1 Layer LSTM



2 Layer LSTM

B) EXEMPLO USANDO CAMADA ACHATADA

```
imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)

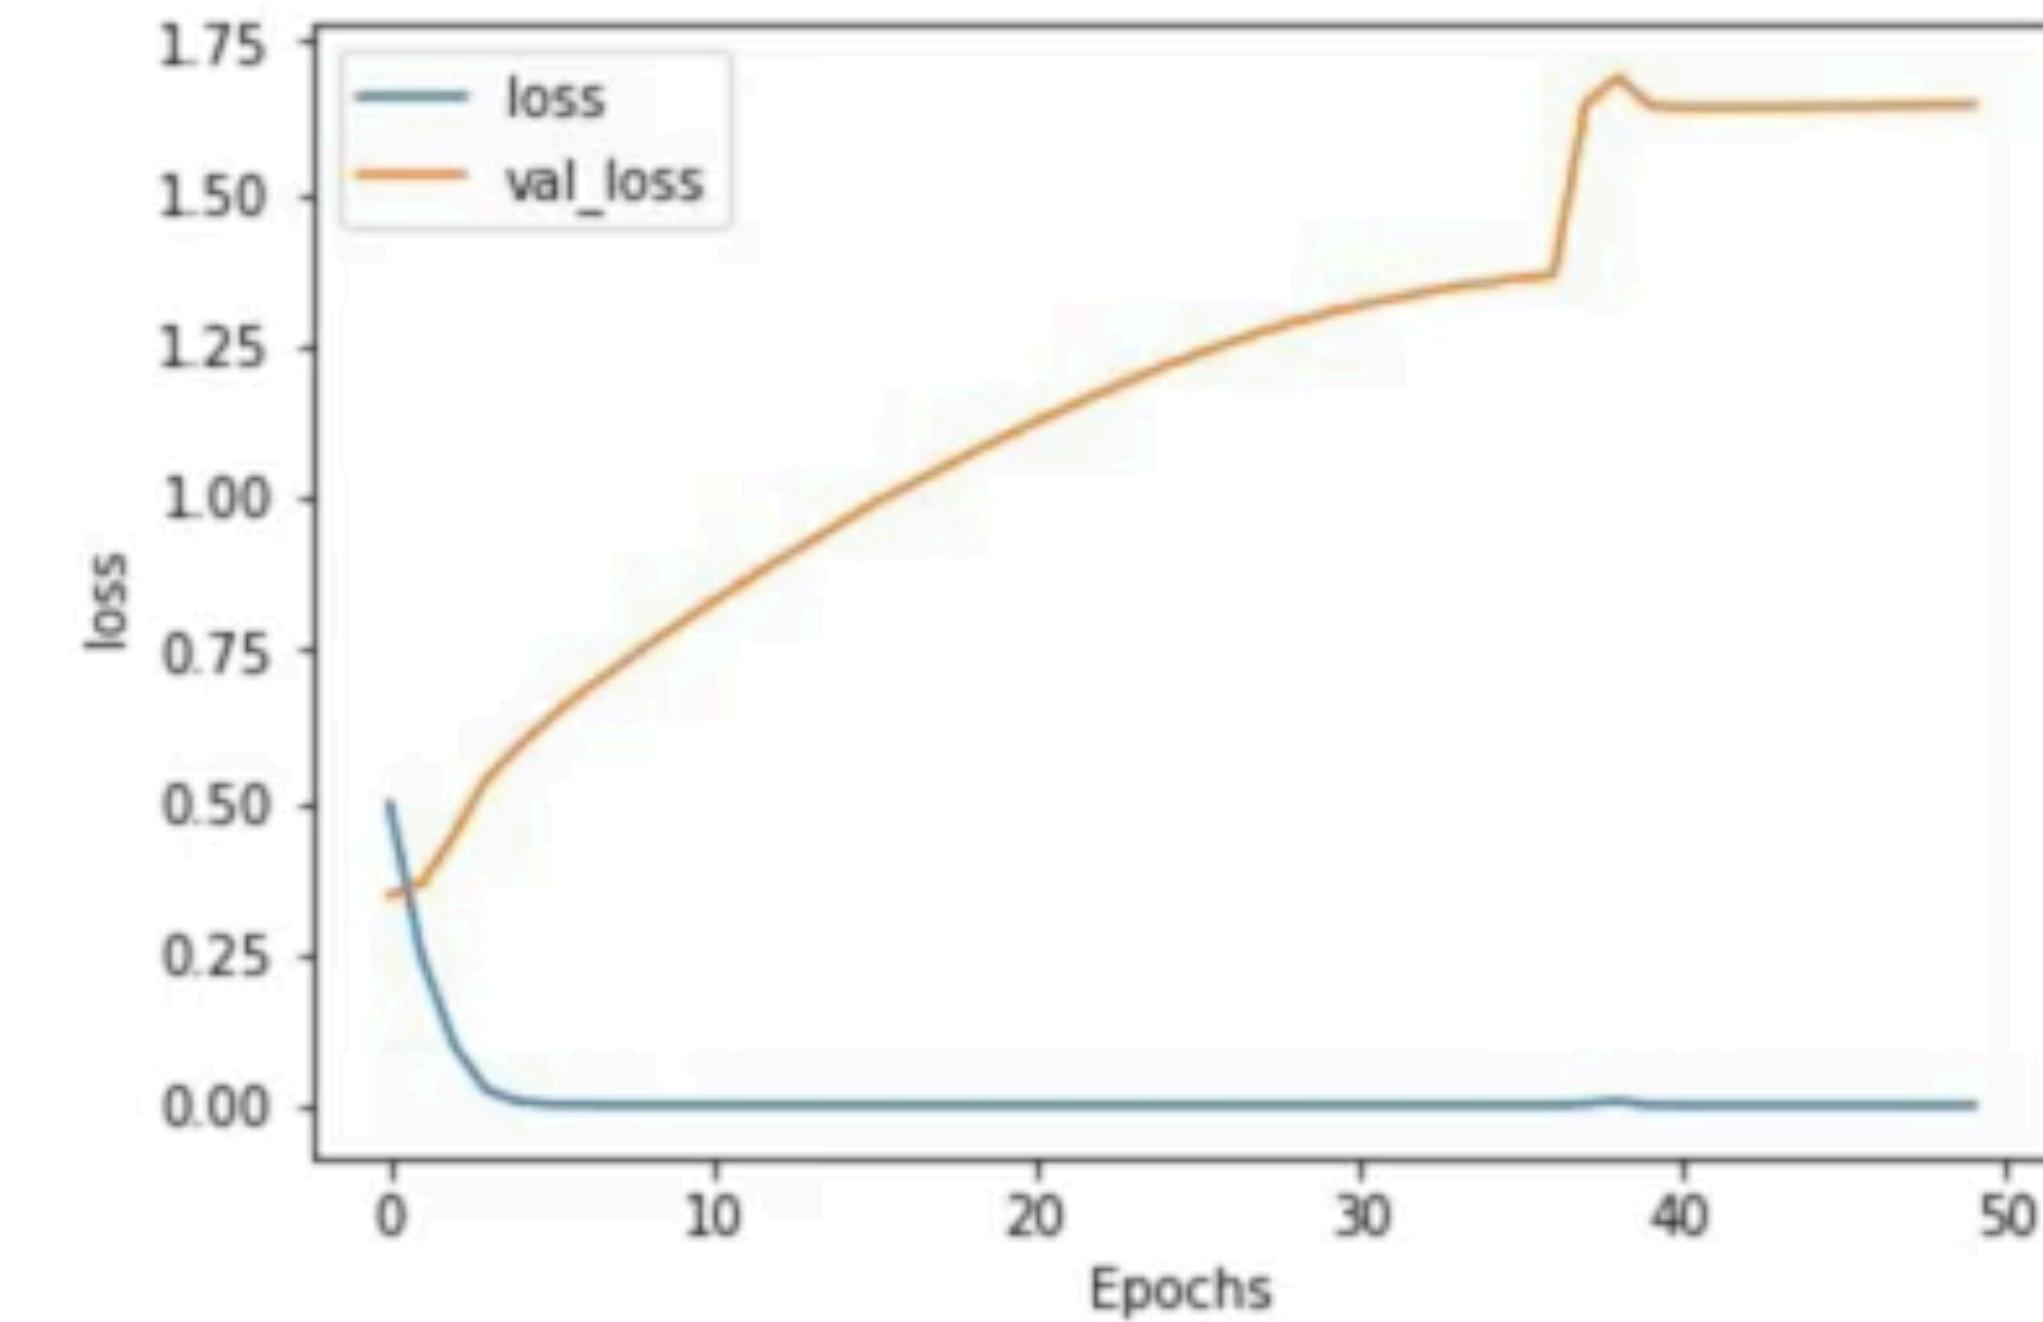
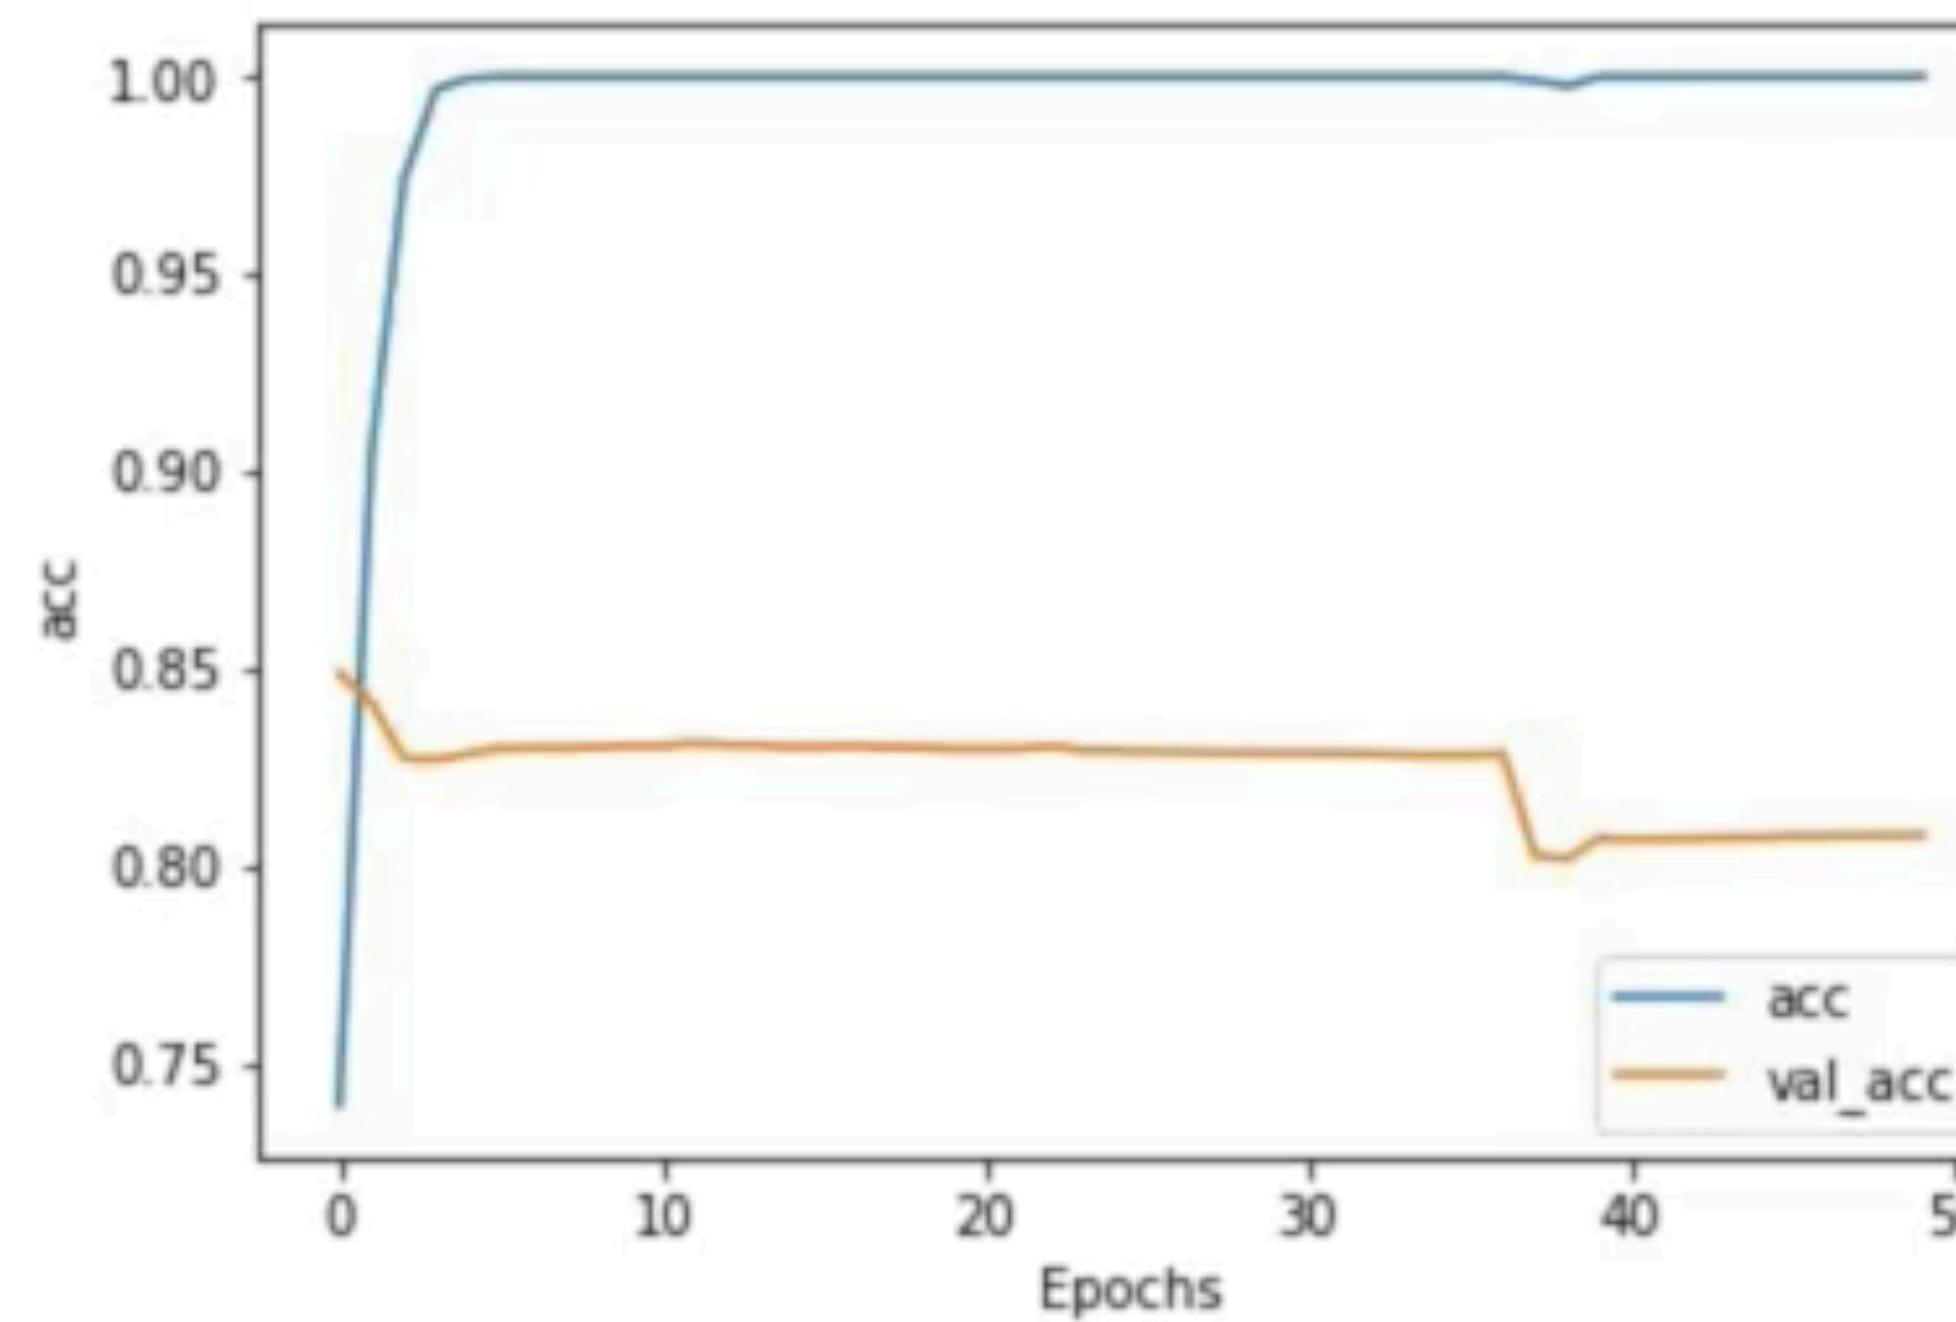
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(6, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=[ 'accuracy' ])

model.summary()
```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 120, 16)	160000
flatten (Flatten)	(None, 1920)	0
dense (Dense)	(None, 6)	11526
dense_1 (Dense)	(None, 1)	7
<hr/>		
Total params: 171,533		
Trainable params: 171,533		
Non-trainable params: 0		





IMDB with Embedding-only : ~ 5s per epoch



 Compartilhar

```
imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)

# Model Definition with LSTM
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),
    tf.keras.layers.Dense(6, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
model.summary()
```

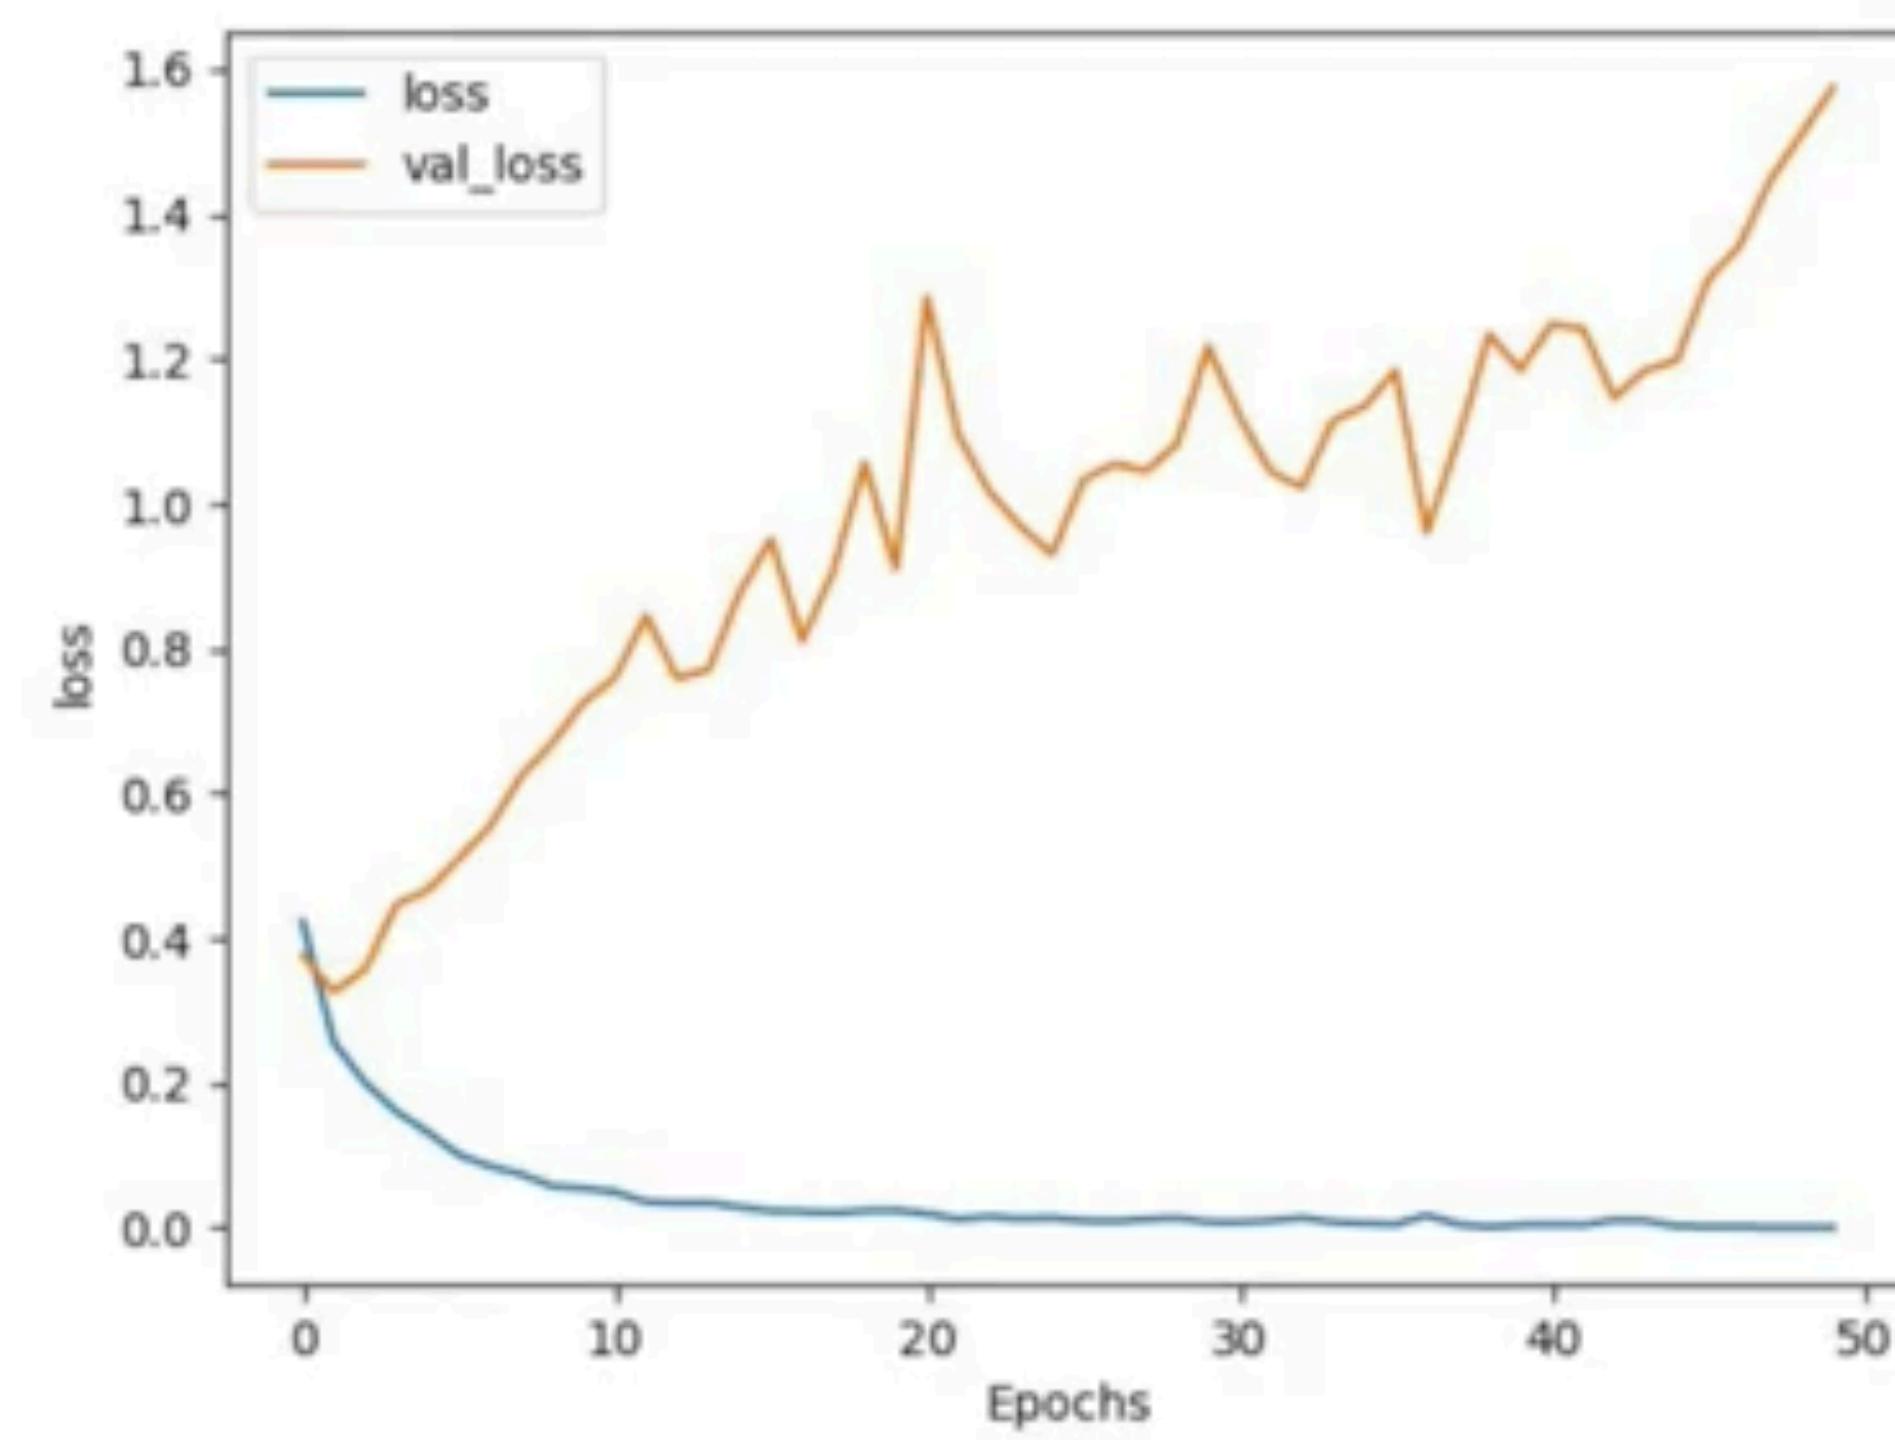
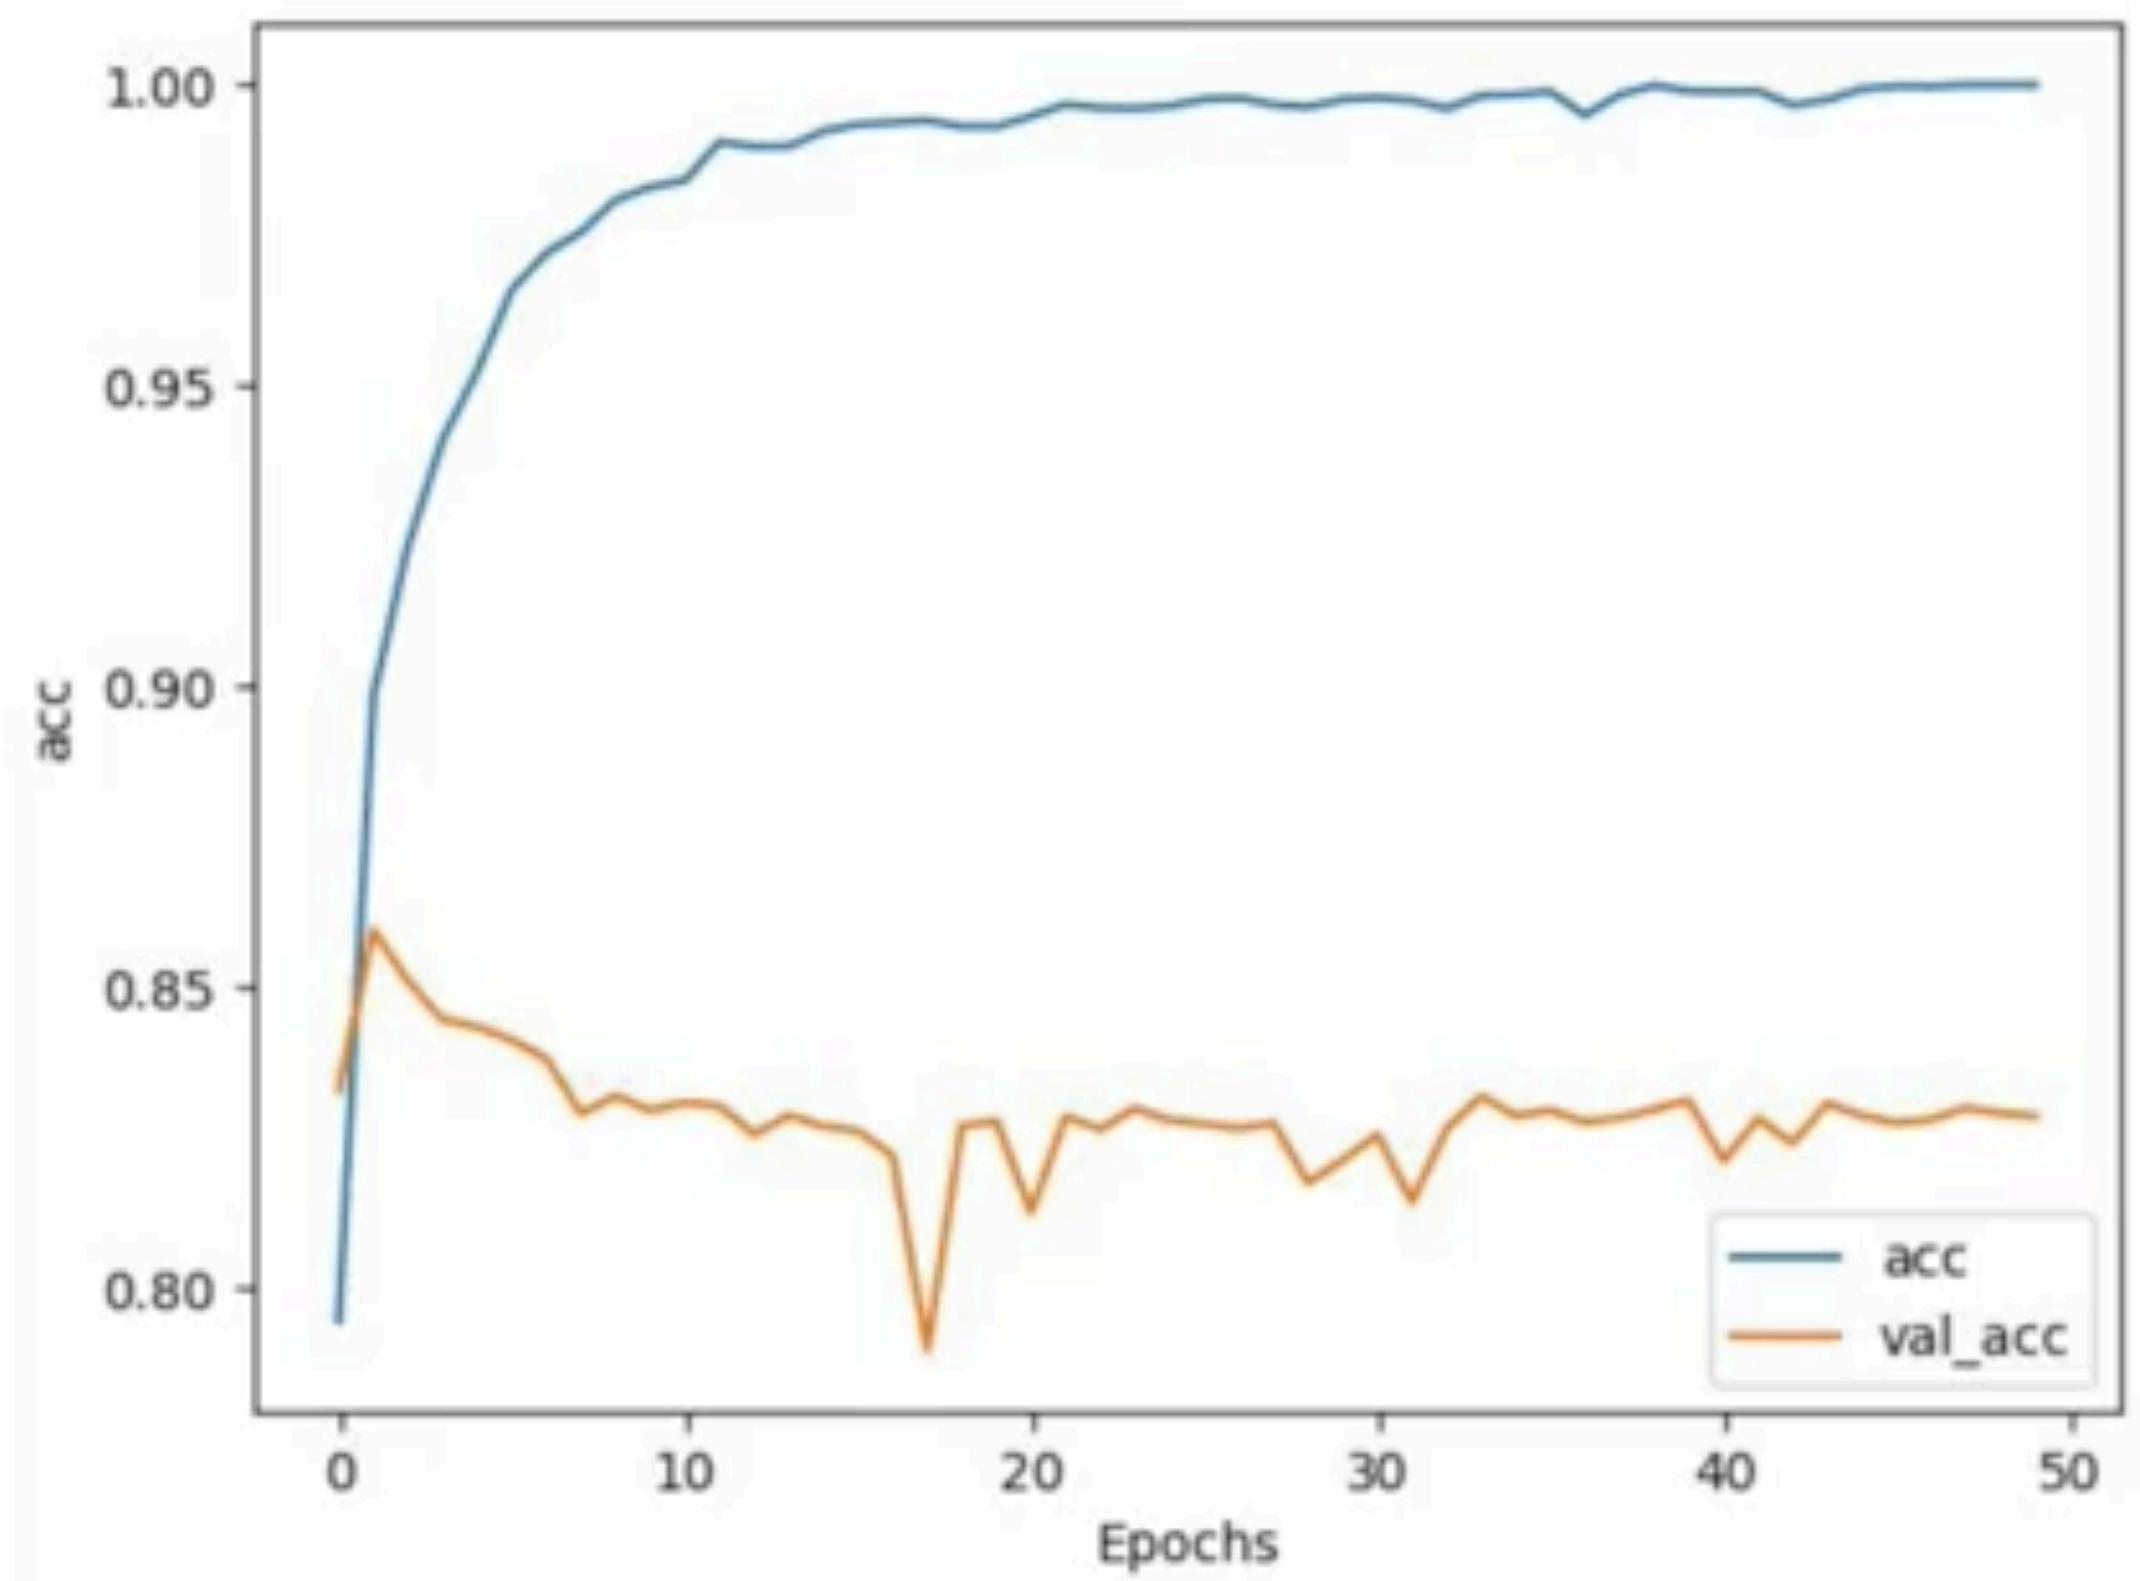
Layer (type)	Output Shape	Param #
embedding_7 (Embedding)	(None, 120, 16)	16000
bidirectional_7 (Bidirection)	(None, 64)	12544
dense_14 (Dense)	(None, 24)	1560
dense_15 (Dense)	(None, 1)	25

Total params: 30,129

Trainable params: 30,129

Non-trainable params: 0





IMDB with LSTM ~43s per epoch



```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.Bidirectional(tf.keras.layers.GRU(32)),
    tf.keras.layers.Dense(6, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=[ 'accuracy'])

model.summary()
```

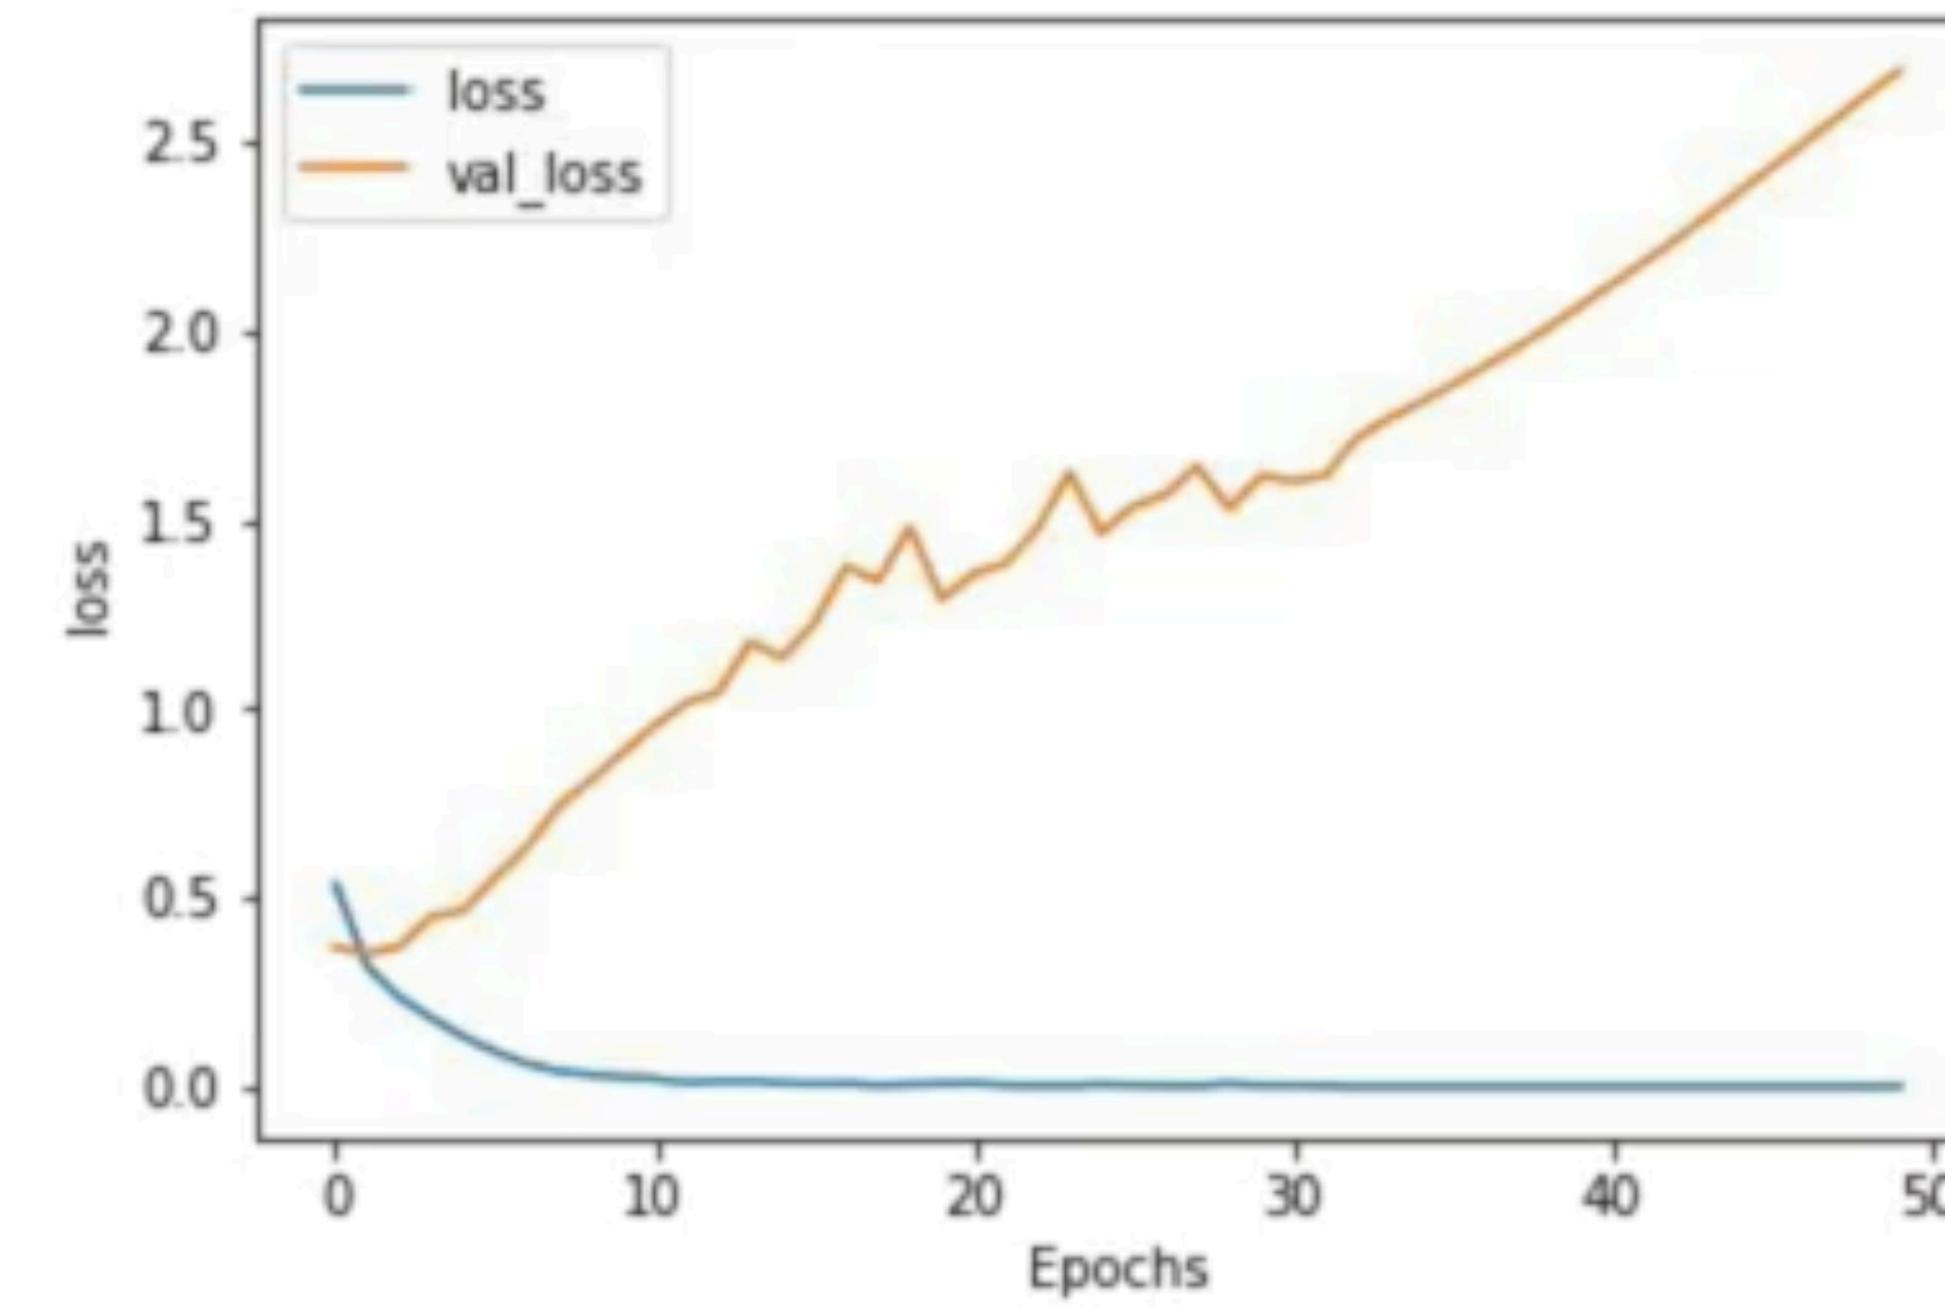
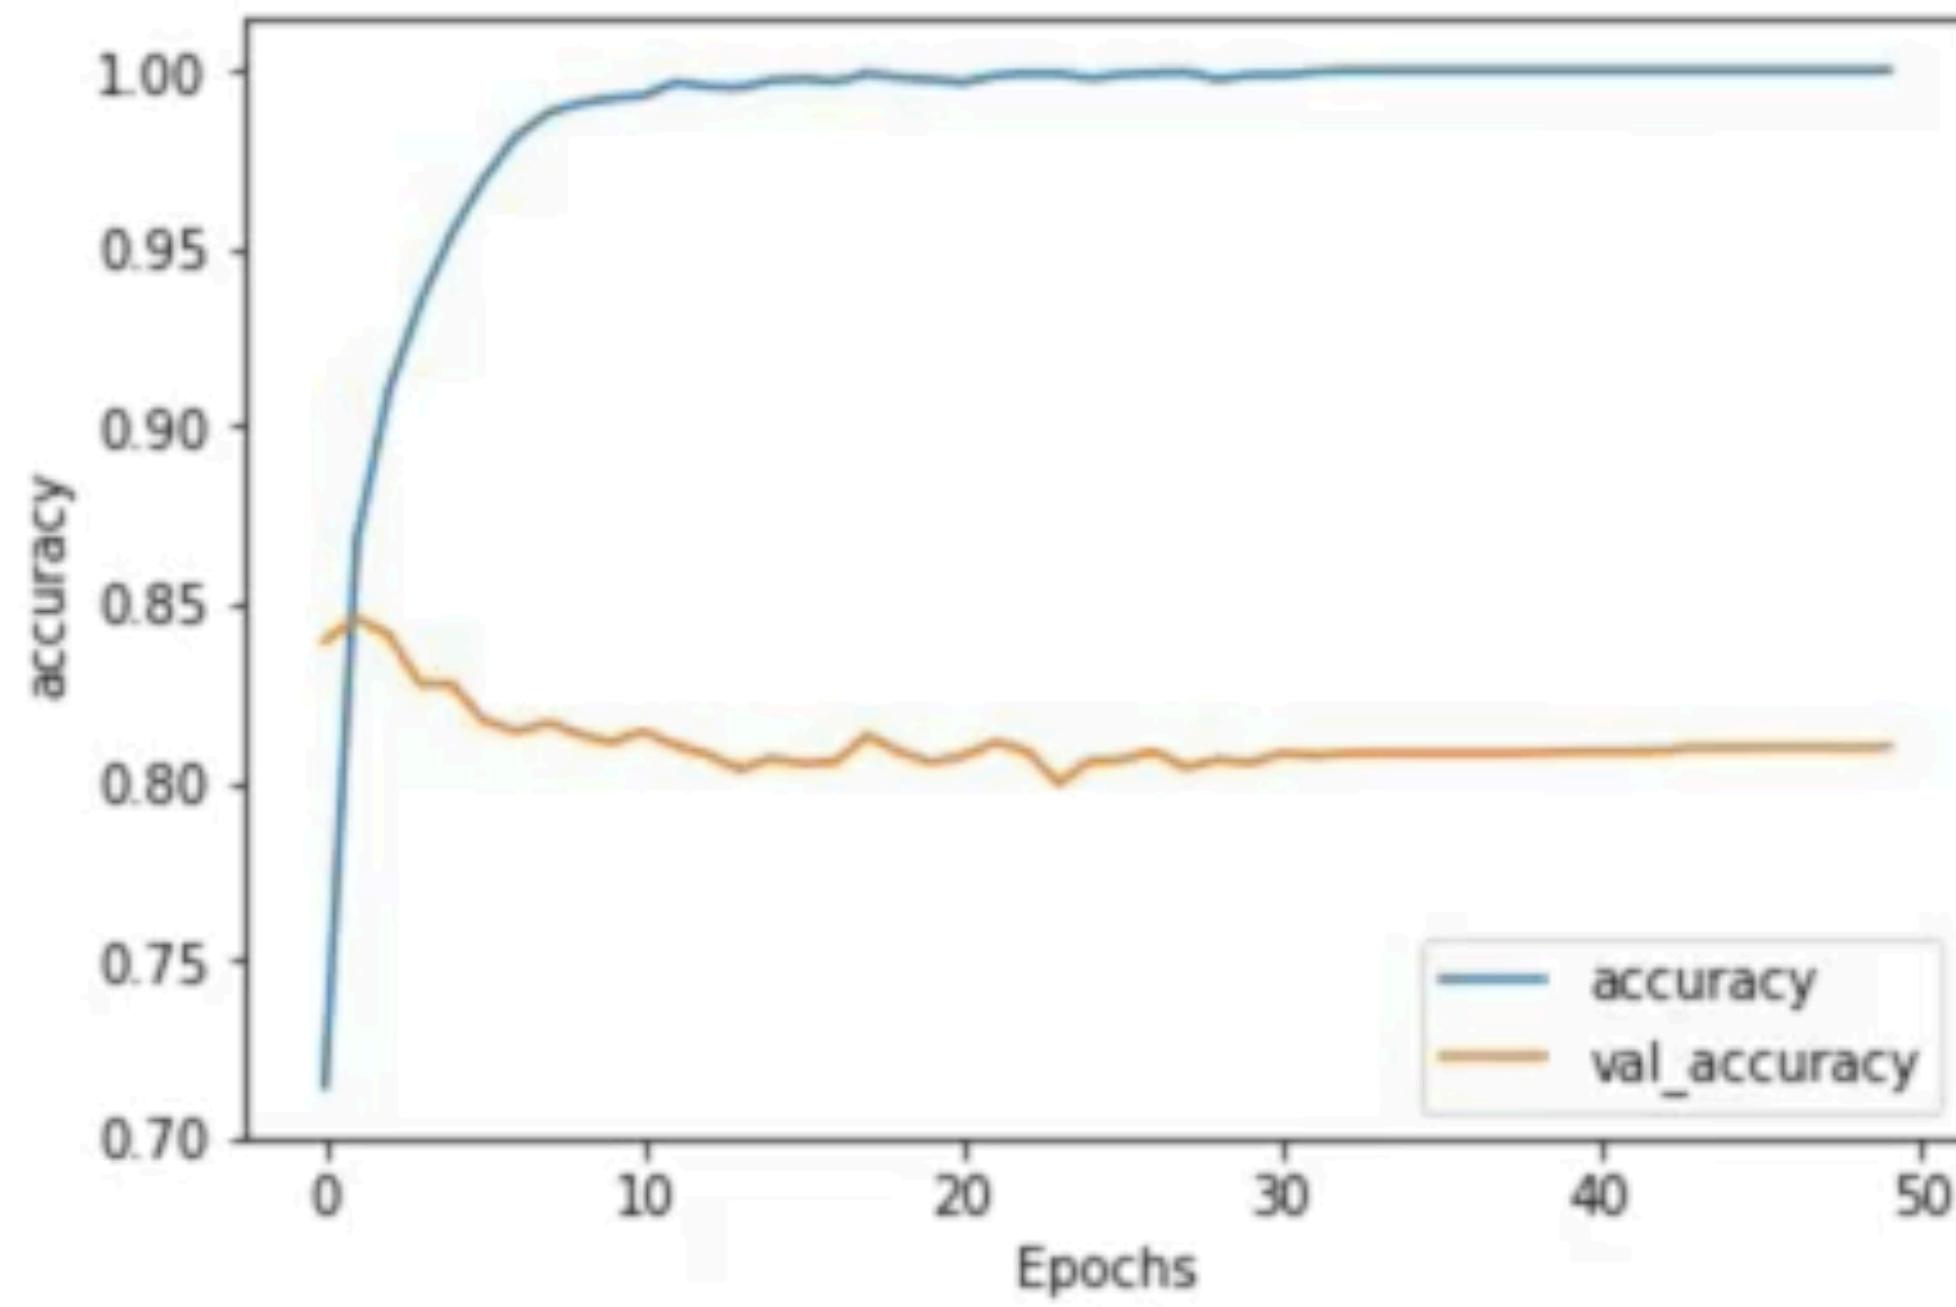
Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 120, 16)	160000
bidirectional_1 (Bidirection)	(None, 64)	9600
dense_2 (Dense)	(None, 6)	390
dense_3 (Dense)	(None, 1)	7

Total params: 169,997

Trainable params: 169,997

Non-trainable params: 0





IMDB with GRU : ~ 20s per epoch



```
# Model Definition with Conv1D
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.Conv1D(128, 5, activation='relu'),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(6, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=[ 'accuracy' ])
model.summary()
```

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 120, 16)	160000

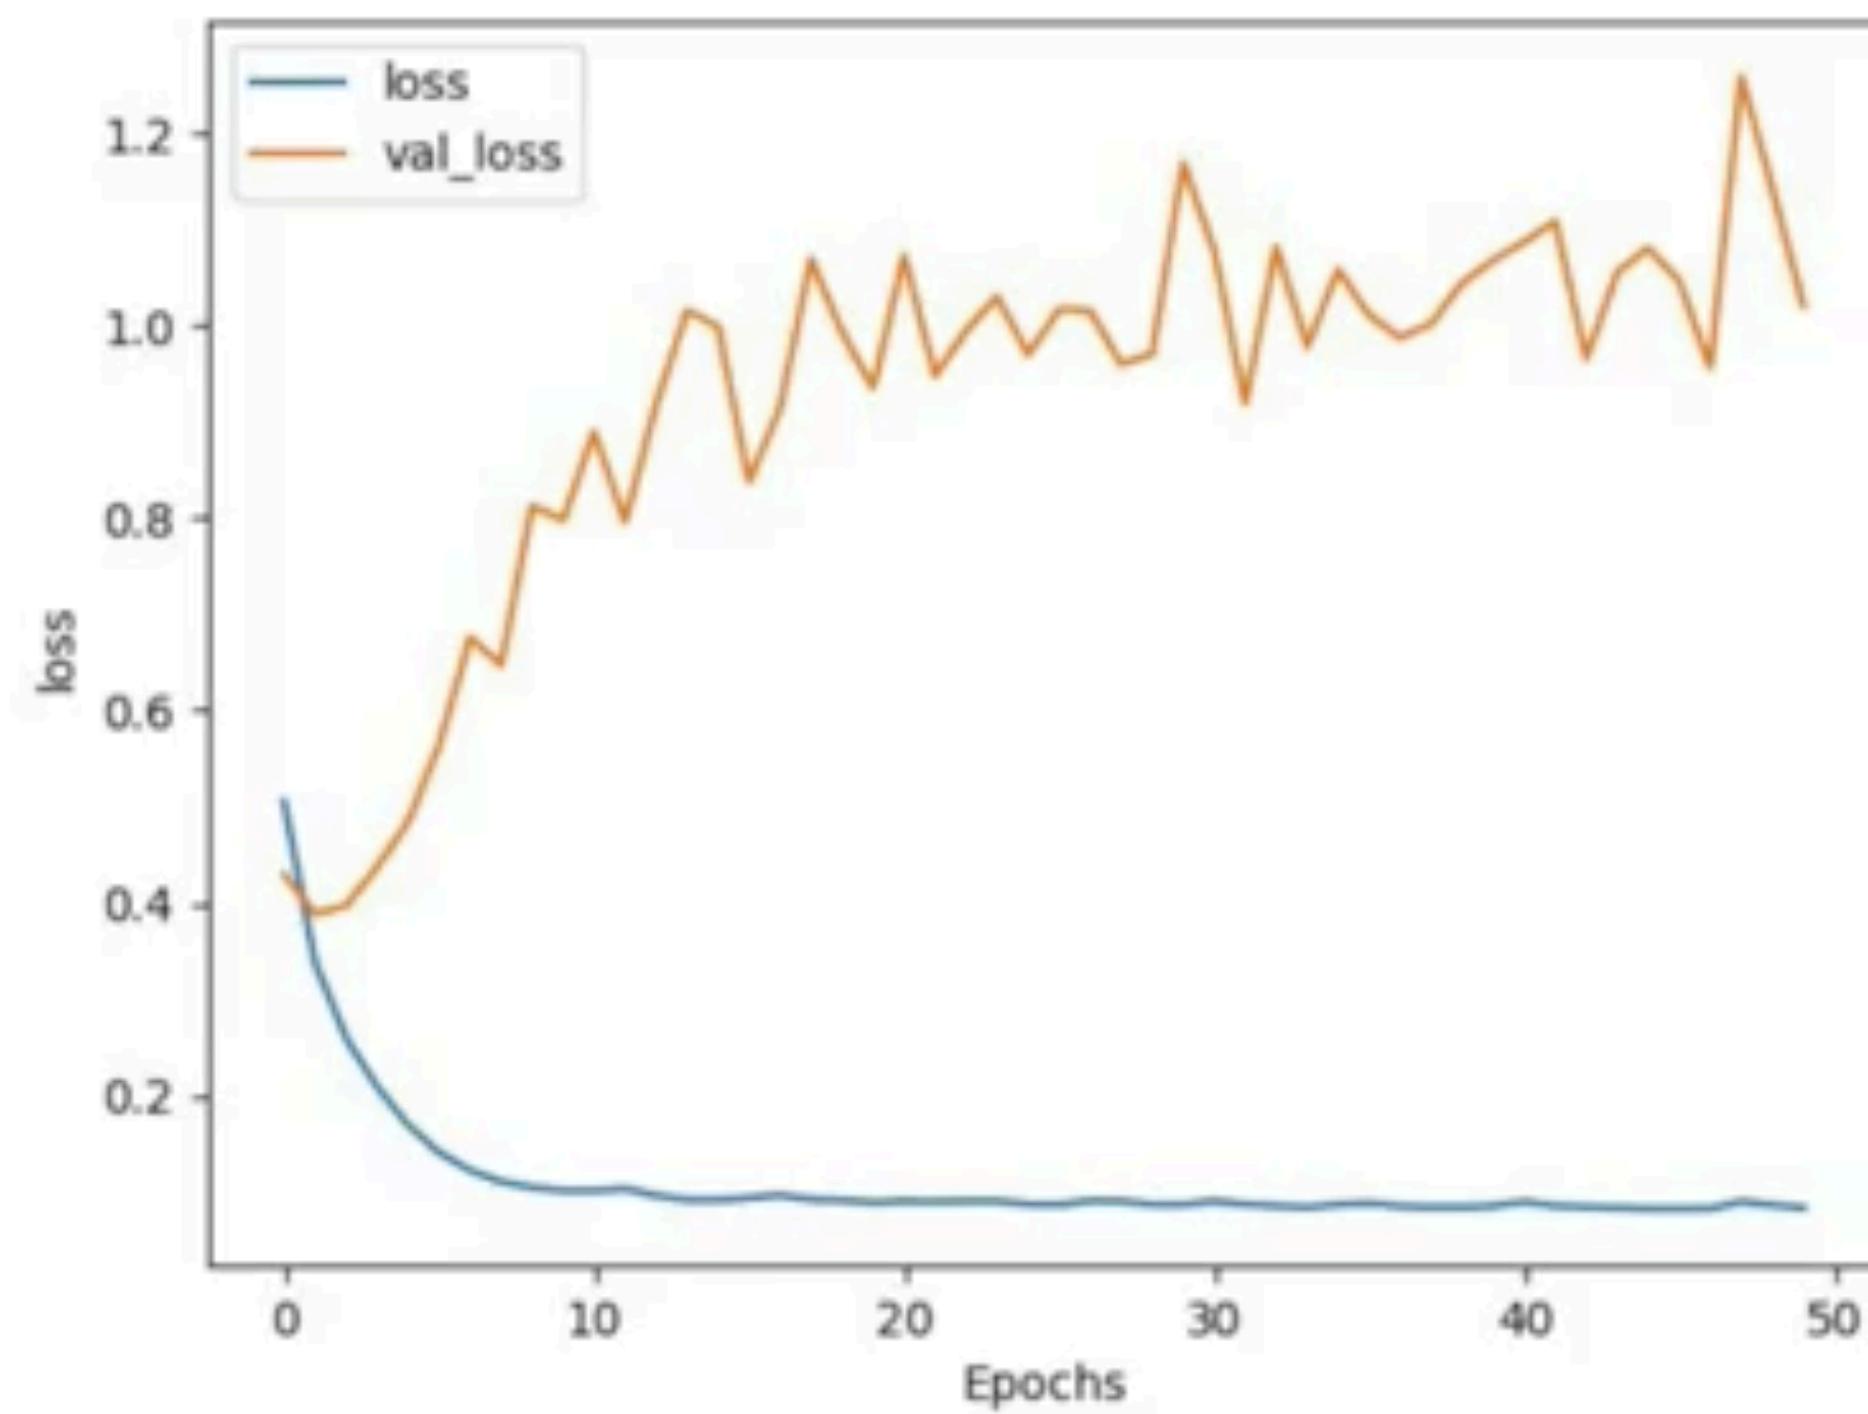
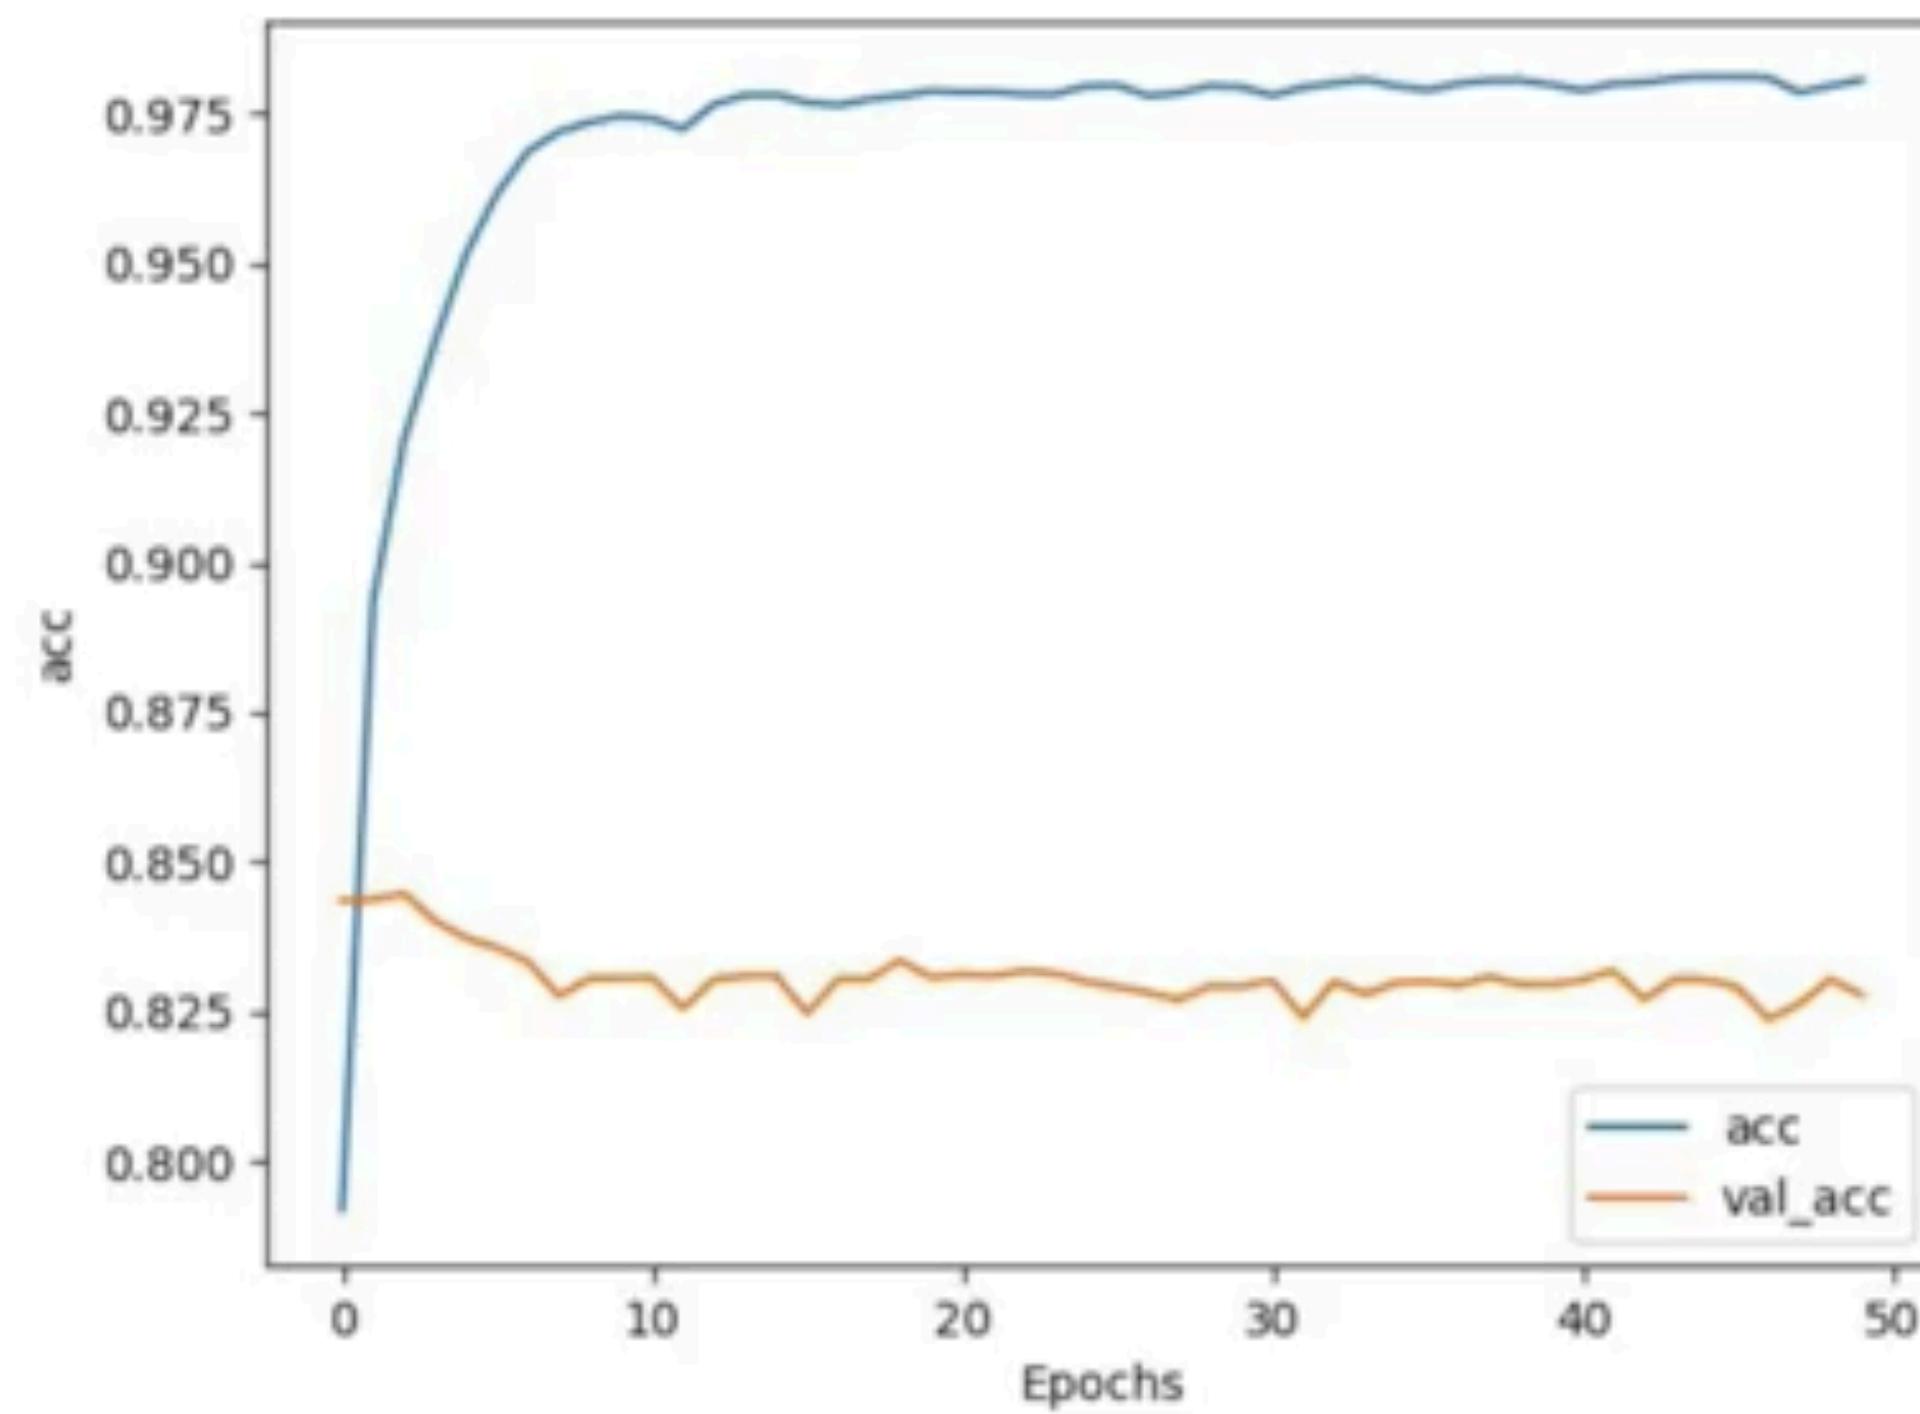
conv1d (Conv1D)	(None, 116, 128)	10368

global_average_pooling1d (G1	(None, 128)	0

dense (Dense)	(None, 6)	774

dense_1 (Dense)	(None, 1)	7
=====		
Total params:	171,149	
Trainable params:	171,149	
Non-trainable params:	0	





IMDB with CNN : ~ 6s per epoch



deeplearning.ai

LINKS ÚTEIS

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%203%20-%20NLP/Course%203%20-%20Week%203%20-%20Lesson%201a.ipynb>

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%203%20-%20NLP/Course%203%20-%20Week%203%20-%20Lesson%201b.ipynb>

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%203%20-%20NLP/Course%203%20-%20Week%203%20-%20Lesson%202d.ipynb#scrollTo=nHGYuU4jPYaj>

EXERCÍCIOS

Explore os modelos criados abaixo:

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%203%20-%20NLP/Course%203%20-%20Week%203%20-%20Lesson%202.ipynb#scrollTo=g9DC6dmLF8DC>

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%203%20-%20NLP/Course%203%20-%20Week%203%20-%20Lesson%202c.ipynb#scrollTo=g9DC6dmLF8DC>

DÚVIDAS?

RAFAELVC2@GMAIL.COM

